

Computer Networks

COL 334/672


Intra-domain Routing

Slides adapted from KR

Sem 1, 2025-26

Intra-domain Routing

Graph Abstraction



graph: $G = (N, E)$

N : set of routers = $\{u, v, w, x, y, z\}$

E : set of links = $\{(u, v), (u, x), (v, x), (v, w), (x, w), (x, y), (w, y), (w, z), (y, z)\}$

$c_{a,b}$: cost of *direct* link connecting a and b
e.g., $c_{w,z} = 5, c_{u,z} = \infty$

How to determine shortest path from one node to all other nodes in a graph?

→ Bellman Ford

Floyd Warshall

→ Dijkstra

- graph is dynamic
 - ↳ node can come & go
 - ↳ edge can " " " "
 - ↳ edge weight can change
- Algorithm should be simple & efficient
- Scalable ⇒ Decentralizes


Distance vector algorithm

Based on *Bellman-Ford* (BF) equation (dynamic programming):

$$D_x(y) = \min_v \{ c(x,v) + D_v(y) \} \leftarrow$$

v: all neighbors of x

Suppose that u 's neighboring nodes, x, v, w , know that for destination z :



$$\begin{aligned} & \min \{ c(u,v) + D_v(z), \\ & \quad c(u,x) + D_x(z), \\ & \quad c(u,w) + D_w(z) \} \\ &= \min \left\{ \begin{array}{l} 2 + 5, \\ 1 + 3, \\ 5 + 3 \end{array} \right\} = 4 \end{aligned}$$

Distance vector algorithm


key idea:

- from time-to-time, each node sends its own distance vector estimate to neighbors
- when x receives new DV estimate from any neighbor, it updates its own DV using B-F equation:

$$D_x(y) \leftarrow \min_v \{c_{x,v} + D_v(y)\} \text{ for each node } y \in N$$

- under minor, natural conditions, the estimate $D_x(y)$ converge to the actual least cost $d_x(y)$

Distance vector: Example



DST	COST	NXT HOP
A	0	A
B	1	B
C	1	C
D	2	C
E	1	E
F	2	E

It #1

A sends : $(A, 0), (B, 1), (C, 1), (E, 1)$

Hears : B : $(B, 0), (A, 1), (C, 1)$

C : $(C, 0), (A, 1), (D, 1), (B, 1)$

E : $(E, 0), (A, 1), (F, 1)$

A sends : $(A, 0), (B, 1), (C, 1), (E, 1), (D, 2), (F, 2)$

Hears : B : $(B, 0), (A, 1), (C, 1), (E, 2), (D, 2)$

C : $(C, 0), (A, 1), (B, 1), (D, 1), (F, 2), (E, 2)$

E : $(E, 0), (A, 1), (B, 2), (F, 1), (D, 2), (C, 2)$

⋮

Good News Travels Fast, Bad News Travels Slow!

- Assume a new node F comes up in the network
- How long does it take for the A to update their routing table?

A — B — C — D — E — F

- In general, good news spreads at rate of one hop per exchange
 - What is the maximum time it can take?
- What happens in case of link failure?

Good News Travels Fast, Bad News Travels Slow!

- Assume the A-B link goes down
- How does the routing table gets updated at B?
- At other nodes?

A ~~X~~ B - C - D - E

Distance to A

∞	2	3	4
----------	---	---	---

3	2	3	4
---	---	---	---

3	4	3	4
---	---	---	---

5	4	5	4
---	---	---	---

5	6	5	6
---	---	---	---

⋮

Another case: if B announces to C false, c will still think it has a better route via D

A ~~X~~ B - C - D - E

∞	2	3	4
----------	---	---	---

∞	4	3	4
----------	---	---	---


5	4	5	4
---	---	---	---

5	6	5	6
---	---	---	---

⋮

How to Handle *Count to Infinity* Problem?

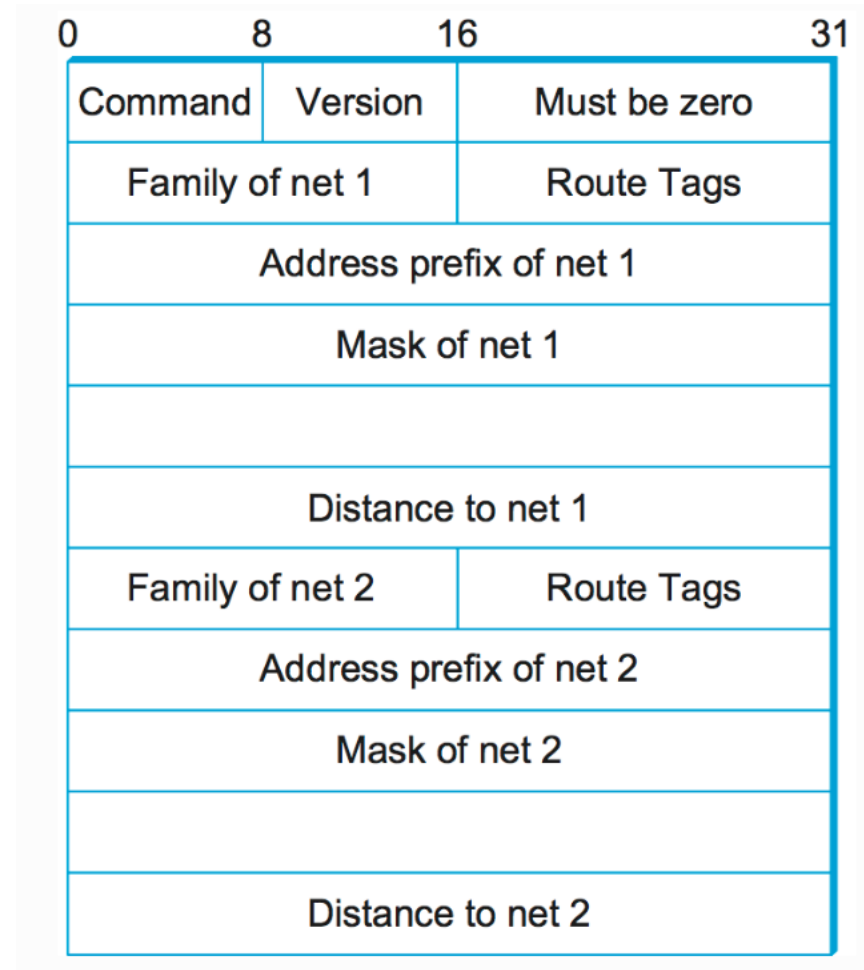
- Split horizon
 - Don't announce the route back to the next hop
 - Does it always work?



How to Handle *Count to Infinity* Problem?

- Split horizon
 - Don't announce the route back to the next hop
- Make infinity smaller
 - Routing Information Protocol (RIP): Based on hop count
 - Used a maximum length of 15
 - **Limitation:** Doesn't work for cases when actual path length is greater than 15


RIP Message Format



RIPv2 Protocol Message Format

Link State Routing Algorithm

- Each node knows the entire network topology including link costs:
 - accomplished via “link state broadcast”
 - all nodes have same info
- Link state broadcast:
 - (x, {list of neighbors, cost}, seq #, TTL)
 - reliable broadcast using acknowledgements
- Once a node has the entire topology:
 - Use **Dijkstra's algorithm** to find the shortest path



Seq # : assess the liveness of msg
→ only forward LSBs of A if
seq # > cur seq # (A)

Dijkstra's link-state routing algorithm

1 *Initialization:*

2 $N' = \{u\}$ /* compute least cost path from u to all other nodes */

3 for all nodes v

4 if v adjacent to u /* u initially knows direct-path-cost only to direct neighbors */

5 then $D(v) = c_{u,v}$ /* but may not be *minimum* cost! */

6 else $D(v) = \infty$

7



8 *Loop*

9 find w not in N' such that $D(w)$ is a minimum

10 add w to N'

11 update $D(v)$ for all v adjacent to w and not in N' :

12 **$D(v) = \min (D(v), D(w) + c_{w,v})$**

13 /* new least-path-cost to v is either old least-cost-path to v or known

14 least-cost-path to w plus direct-cost from w to v */

15 *until all nodes in N'*


Dijkstra's algorithm: an example

LSA: periodic
Triggered &
updates in
N/w

Step	N'	^V D(v),p(v)	^W D(w),p(w)	^X D(x),p(x)	^Y D(y),p(y)	^Z D(z),p(z)
0	u	2, u	5, u	1, u	∞	∞
1	u, x	2, u	4, x	-	2, x	∞
2	u, x, v	2, u	4, x	-	2, x	∞
3	u, x, v, y	-	4, x	-	2, x	4, y
4	u, x, v, y, w	-	-	-	-	4, y
5	u, x, v, y, w, z	-	-	-	-	-

After adding node a to N'
& neighbors of a not in N'

$$D(b) = \min(D(b), D(a) + e(a, b))$$



Two popular link state routing protocols: Intermediate System to Intermediate System (IS-IS), Open Shortest Path First (OSPF)

Setting Link Weights

- How to set link weights?
 - Static: 1, bandwidth of link, cost of the link
 - Adaptive: load on the network
- How to set weights based on load on the network?
 - Send ECHO packets to measure RTT
 - RTT depends on queuing delays (load on the network)
- Is it a good strategy?
 - Can lead to path oscillations, routing loops
 - Variations in end-to-end delay and packet reordering
- Link weight setting is an art!
 - Knobs for traffic engineering
 - Reversal of cause and effect