

# Computer Networks

## COL 334/672

Congestion Control


*Slides adapted from KR*

Sem 1, 2025-26

# Quiz on Moodle

Password: reliability

# Recap: Congestion Control



$$r_1 + r_2 \leq C$$

## Approach for designing CCA

### ① N/w + end hosts

Routers signal end host to slow down when congestion is happening

CONS: Increases complexity


END-TO-END PRINCIPLE

### End-host only

Each end host independently decide the rate w/o any assistance from the n/w

Adapt send rate such that n/w does not get congested

# Metrics for evaluating a Congestion Control Algorithm



- ① High goodput
- ② Low latency
- ③ Fairness

Across a diverse set of n/w conditions

# Designing an End-to-end Congestion Control Algorithm



Probe the network by increasing the sending rate.

Show down if you detect congestion

Decrease  
sending rate

Two Fundamental Questions

- ①. how to detect congestion
- ②. how much to increase/decrease

# How to increase/decrease the sending rate?

- Many functions are possible
- *Class of linear* control algorithms:
  - $w_{i+1} = aw_i + b$   $\rightarrow$   $c$
  - $a$  and  $b$  are different in case of increase and decrease
- Additive Increase, Multiplicative Decrease (AIMD)
  - $w_{i+1} = w_i + b$   $\rightarrow$  increase function  $b > 0$
  - $w_{i+1} = \underbrace{aw_i}_{a < 1}$   $\rightarrow$  decrease function

→ Assuming rate is controlled using congestion window (cwnd)

# Congestion control: AIMD

- *approach*: senders can increase sending rate until congestion occurs, then decrease sending rate on congestion

$$w_{i+1} \leftarrow w_i + \text{MSS}(\text{every time unit})$$


## Additive Increase

increase sending rate by 1 maximum segment size (MSS) every time unit until congestion detected

$$w_{i+1} \leftarrow w_i / 2$$

## Multiplicative Decrease

cut sending rate in half at each congestion event



→ **AIMD** sawtooth behavior: *probing* for bandwidth

# Congestion control: AIMD

- *approach*: senders can increase sending rate until congestion occurs, then decrease sending rate on congestion

## Additive Increase

increase sending rate by 1 maximum segment size (MSS) every **time unit** until **congestion detected**

## Multiplicative Decrease

cut sending rate in half at each congestion event

Questions?

- How to detect congestion?
- Why AIMD?
- What is the time unit?

Two signals of congestion

① Increase in RTTs  $\Rightarrow$  queuing delay is increasing  $\Rightarrow$  congestion

② Packet loss

①  $\downarrow$  ACK Timeout


② Triple duplicate ACK



# Why AIMD? $\rightarrow$ TCP fairness [MIMD, MIAD, AIAD]

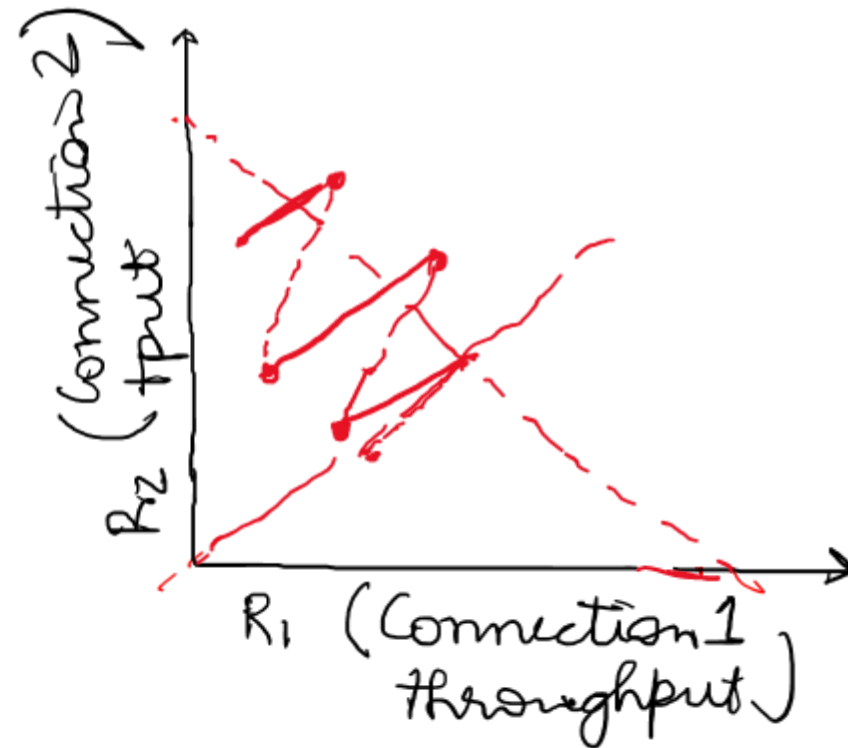
**TCP fairness:** if  $K$  TCP sessions share same bottleneck link of bandwidth  $R$ , each should have a long-term average rate of  $R/K$

$R/2$



$$R_1 + R_2 \leq R$$

AIMD is TCP fair



# Window Update in AIMD

- **Goal:** For additive increase, update window size by 1 MSS every RTT. But when?

- Window updated as ACKs arrive


- How much should be the update?

$\boxed{cwnd} \pm 1 \text{ MSS}$  every RTT or every 1 second

# of unacked bytes that TCP has

A

B



$cwnd$  (in bytes)

How many segments are in flight?:

$$\boxed{cwnd / MSS}$$

want to increase 1 MSS every RTT

$$\left[ cwnd \leftarrow cwnd + \frac{1 \text{ MSS} \times \text{MSS}}{cwnd} \right]$$

$$\frac{cwnd}{RTT} \text{ bytes per second} \rightarrow \frac{cwnd + MSS}{RTT}$$

[Why not increase ~~stand~~ by 1 every ACK]

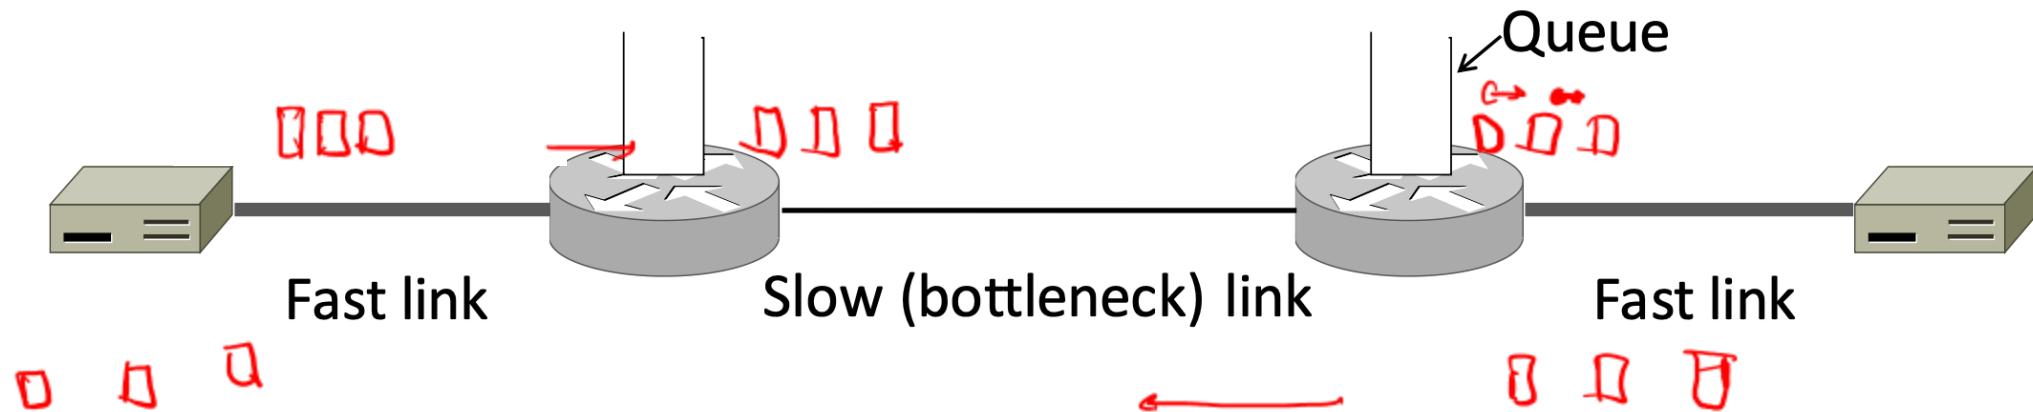
↓  
Exponential increase

- Because TCP updates based on ACKs, TCP is self-clocking

Benefit: Self-clocking smooths packet sending rate

# TCP Self-Clocking

- ACK clocking smooths packet sending rate



# In the beginning of the connection...

- What *cwnd* size should we start with?
  - **Goal:** We want to quickly near the *right* rate.
  - A linear increase with a small value of *cwnd* is painfully slow!
- What are the options?
  - Start with a large value of *cwnd*
  - Start with a small value of *cwnd* but increase it faster

$IW = 2 \rightarrow \boxed{IW = 10}$

↓  
exponential increase

instead of changing *cwnd* by 1MSS every RTT

double the *cwnd* every RTT