

Computer Networks

COL 334/672

Congestion Control

Slides adapted from KR

Sem 1, 2025-26

Recap

- Congestion control:
 - Network-assisted vs end-to-end
- Designing an end-to-end congestion control
 - **General approach:** probe the network by increasing the **sending rate**, decrease when you infer congestion
 - **Additive Increase, Multiplicative Decrease:** TCP fair
 - **Inferring congestion:** use packet loss as an indicator of congestion / other indicators: delays

sending rate controlled
using sliding window
mechanism


$slos \leq \min(cwnd, rwnd)$

cwnd: congestion window
(internal to sender)

rwnd: flow control
window advertised by
receiver

Honing CCA Further

- **One specific instance of AIMD:** Increase sending rate by 1 MSS every RTT, decrease to half on event of congestion
 - **When to increase:** on reception of every new ACK, since ACK indicates a packet has been received
 - **How much to increase:** $MSS * MSS / cwnd$
- **What to do in the beginning of the connection?**
 - Linear increase is too slow
 - **Q:** Given a network with RTT of 80ms, initial congestion window of 1 MSS. MSS is 1000 bytes. How long it will take to get to an average rate of 10 Mbps?



$$\text{Rate} = \frac{cwnd}{RTT}$$

$$cwnd = 100 \text{ MSS}$$

99 RTTs to increase
to 100 MSS if $IN = 1 \text{ MSS}$

$$99 \times 0.08 \approx 8 \text{ seconds}$$

What to do in the beginning..

- How to ramp up faster?

OPT-1

Larger initial window (say 100 MEG).

⇒ But it can overwhelm the slower networks

OPT-2


Start with a small window but ramp up faster than linear

= one idea: increase exponentially

TCP slow start


- when connection begins, increase rate exponentially until first loss event:
 - initially **cwnd** = 1 MSS
 - double **cwnd** every RTT
 - done by incrementing **cwnd** for every ACK received
- *summary*: initial rate is slow, but ramps up exponentially fast

Until when?




TCP Slow Start and Congestion Avoidance

- Uses a threshold, *ssthresh*, after which TCP enters **congestion avoidance**
- *What happens when a loss occurs?*



What happens when a loss occurs?

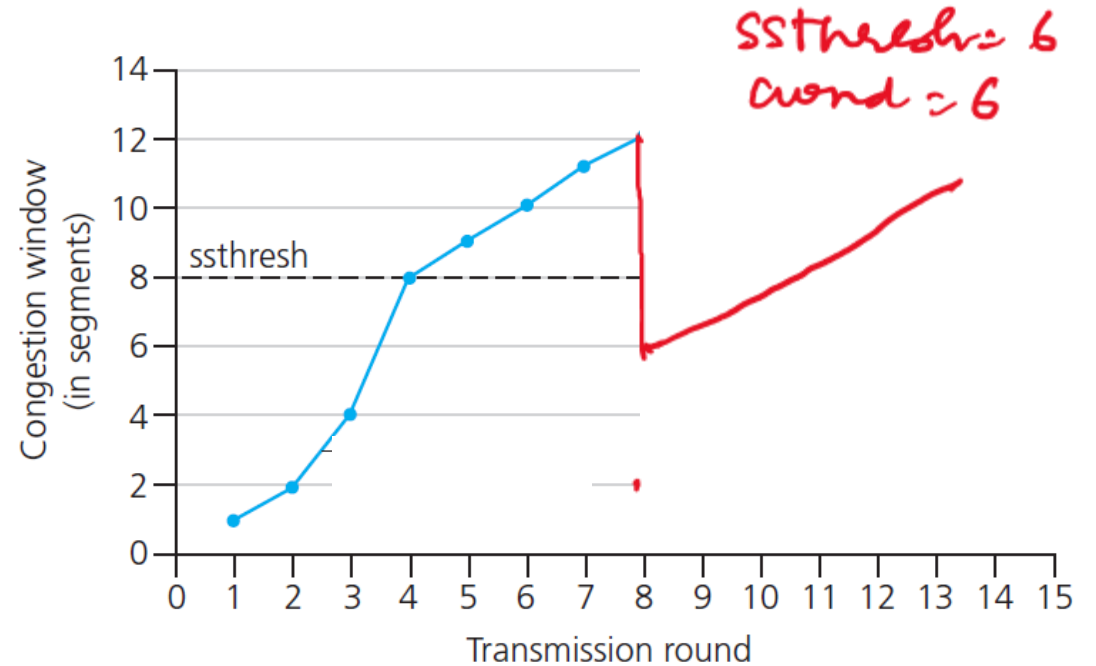
- When loss is due to timeout:
severe congestion!!
 - Set *ssthresh* to *cwnd*/2
 - Reset *cwnd* to 1
 - Enter **slow start phase**



What happens when a loss occurs?

- When loss is due to **triple duplicate ACK**: congestion is not severe!
 - Set *ssthresh* to $cwnd/2$
 - On receipt of another duplicate ACK, send 1 new segment
 - Once a new ACK arrives, set $cwnd = ssthresh$ (or $cwnd/2$)
 - Enter **congestion avoidance** phase

TCP Reno: Fast retransmit, fast recovery!



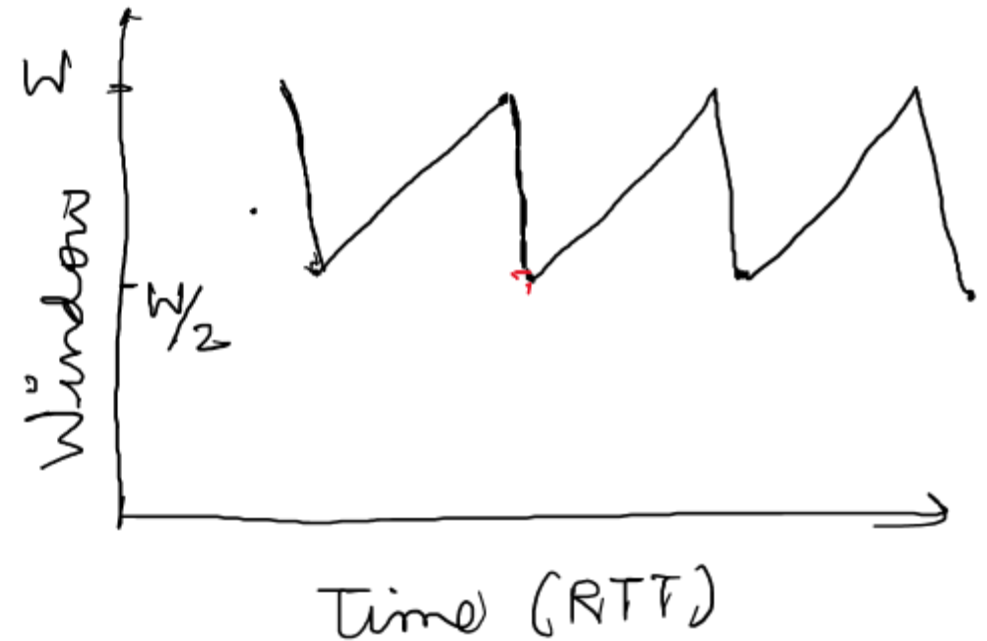
TCP Reno Throughput: Macroscopic Description

- Throughput: area under the curve

$$\text{Average window} = \frac{3W}{4}$$

$$\text{Average Tput} = \frac{3W \times \text{MSS}}{4 \text{RTT}}$$

HW: derive tput if packet loss probability is p



Still slow if B/w is high or RTT is high

i.e. $B/w \times \text{RTT}$ is high

↑

Inefficient for networks with high bandwidth delay product!


Can we do faster?

Is there a better way than AIMD to “probe” for usable bandwidth?

■ Insight/intuition:

- W_{\max} : sending rate at which congestion loss was detected
- congestion state of bottleneck link probably (?) hasn't changed much
- after cutting rate/window in half on loss, initially ramp to to W_{\max} *faster*, but then approach W_{\max} more *slowly* → *Binary search*

TCP BIC : Binary Increase Congestion Control



Improving further

- TCP Reno: Window increases by 1 MSS per RTT
- What happens when there are flow with different RTTs?

RTT unfairness

Flow with smaller RTT will ramp up faster ;

TCP CUBIC

- K: point in time when TCP window size will reach W_{max}
 - K itself is tunable
- increase W as a function of the *cube* of the distance between current time and K

$$W(t) = C(t - K)^3 + W_{max}$$

$$K = \sqrt[3]{\frac{W_{max}\beta}{C}}$$

K is the time it takes to reach W_{max}

$$W(K) = W_{max}$$

$$W(0) = (1 - \beta) W_{max}$$

- TCP CUBIC default in Linux, most popular TCP for popular Web servers

[TCP BBR is replacing CUBIC]