

Computer Networks

COL 334/672


Application Layer

Slides adapted from KR

Sem 1, 2025-26

Internet Layered Architecture

- *application*: supporting network applications
 - HTTP, IMAP, SMTP, DNS
- *transport*: process-process data transfer
 - TCP, UDP
- *network*: routing of datagrams from source to destination
 - IP, routing protocols
- *link*: data transfer between neighboring network elements
 - Ethernet, 802.11 (WiFi), PPP
- *physical*: bits “on the wire”



How to design a networked application?

① What is that application: chat application.

② scale: 100 users (low budget)



For N/w communication: Socket APIs

- UDP socket
- TCP socket → reliability

N/w app architecture



[Intermittent connectivity / May be behind a NAT]

Server

- dedicated machine (always on)
- publicly accessible (Public IP, port)

P2P architecture Peer to peer



(Both are intermittently connected)

↳ Establishing connectivity (Torrent) is challenge

↳ No resources / self-scaling n/w

↳ Reliability / Churn

Networked Application Architecture

- Publish subscribe architecture



What we are going to cover

- HTTP

(2)

- Email

(1)

- DNS

→ Infrastructure Appn

- P2P

(3)

- Video Streaming

(4)

Killer Applications

Web and HTTP

Browser downloads base HTML
which has info about all other
objects

- World Wide Web or Web was the 2nd **killer app** over the Internet
- **Goal of Web**: organize and retrieve information over Internet
- Web is based on two sister protocols: **HTML** and **HTTP**
- **HTML**: HyperText Markup Language
 - language to create and structure web page
 - web page consists of **objects**, each of which can be stored on different Web servers
 - **base HTML-file** which includes several referenced objects, each addressable by a URL

`www.wikipedia.org/wiki/http`

host name

path name

example.com

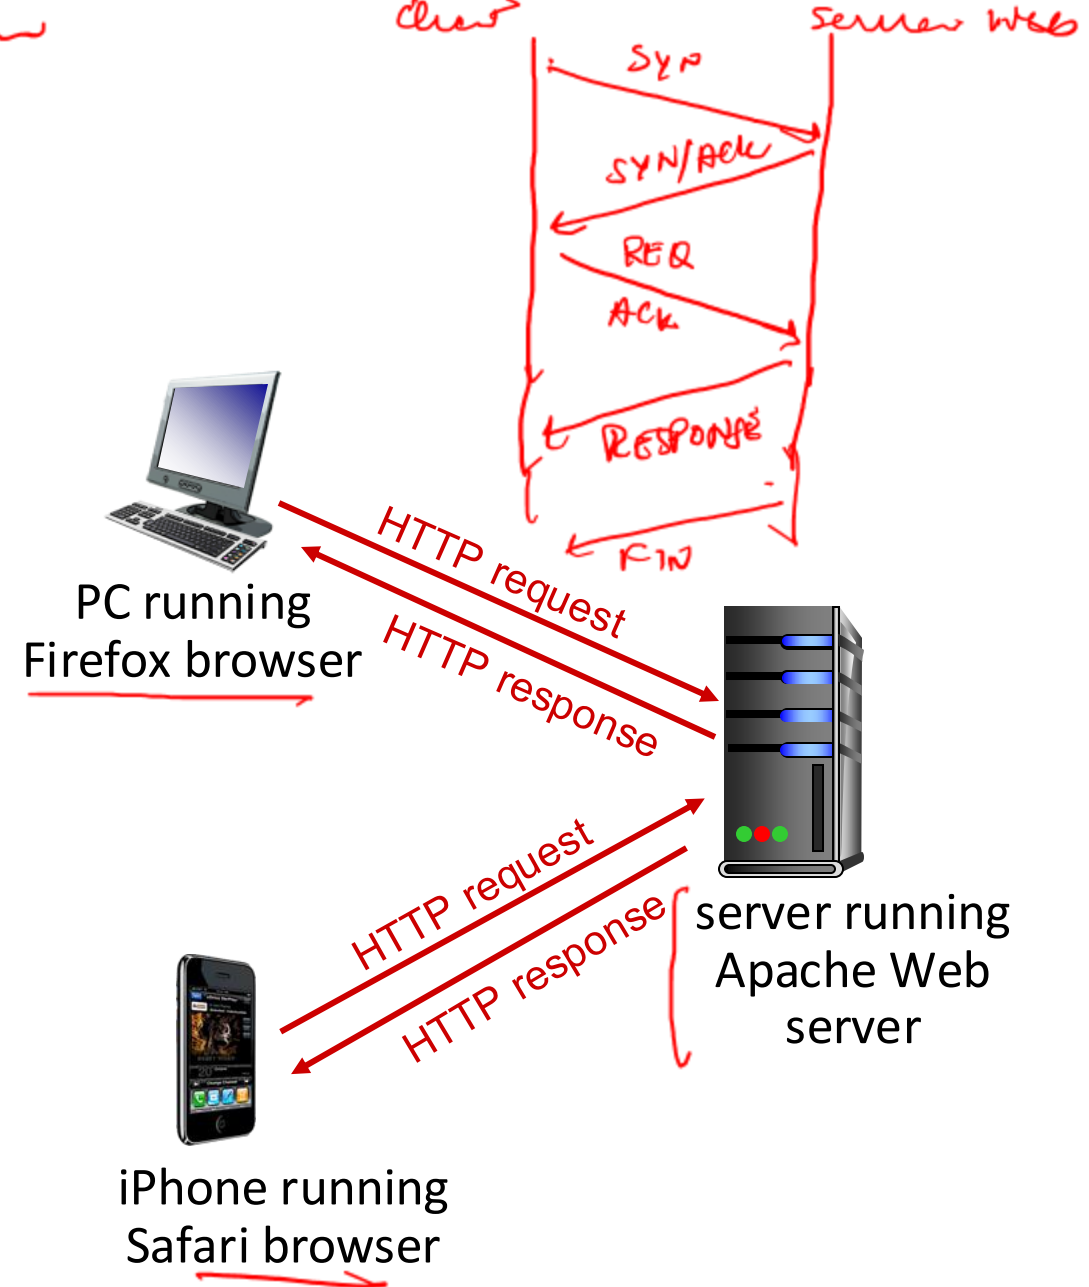
HTTP overview

HTTP: hypertext transfer protocol

- Request/response protocol
- Stateless protocol [scalability, simple]
- Which transport?

TCP! Reliability
↳ Congestion control


2 RTTs per object



Downloading Multiple Objects

10 objects : 20 RTT (Non-persistent connection)

11 RTT (persistent connection)



Request (Keep Alive)

Persistent HTTP (HTTP 1.1)

Non-persistent HTTP issues:

- requires 2 RTTs per object
- OS overhead for *each* TCP connection
- browsers often open multiple parallel TCP connections to fetch referenced objects in parallel

Persistent HTTP (HTTP1.1):

- server leaves connection open after sending response
- subsequent HTTP messages between same client/server sent over open connection
- client sends requests as soon as it encounters a referenced object
- as little as one RTT for all the referenced objects (cutting response time in half)

HTTP

- Overview of HTTP
- **Request/response message format**
- State management
- Caching
- Request pipelining

HTTP request message

cse.iitd.ac.in / ~ cs . . .
↓ Find server (DNS)
IP address
↓ download index.html

- two types of HTTP messages: *request, response*
- HTTP request message:
 - ASCII (human-readable format)

request line (GET, POST,
HEAD commands)

Name, Value
↓
HTTP


header
lines

```
GET /index.html HTTP/1.1\r\n
Host: www-net.cs.umass.edu\r\n
User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X
10.15; rv:80.0) Gecko/20100101 Firefox/80.0 \r\n
Accept: text/html,application/xhtml+xml\r\n
Accept-Language: en-us,en;q=0.5\r\n
Accept-Encoding: gzip,deflate\r\n
Connection: keep-alive\r\n
\r\n
```

carriage return character
line-feed character

carriage return, line feed
at start of line indicates
end of header lines

HTTP request message: general format



Other HTTP request messages

post data

POST method:

1

- web page often includes form input
- user input sent from client to server in entity body of HTTP POST request message

GET method (for sending data to server):

2

- include user data in URL field of HTTP GET request message (following a '?'):

`www.somesite.com/animalsearch?monkeys&banana`

HEAD method:

- requests headers (only) that would be returned *if* specified URL were requested with an HTTP GET method.

2

PUT method:

- uploads new file (object) to server
- completely replaces file that exists at specified URL with content in entity body of POST HTTP request message

HTTP response message

status line (protocol
status code status phrase)

header
lines

HTTP/1.1 200 OK
Date: Tue, 08 Sep 2020 00:53:20 GMT
Server: Apache/2.4.6 (CentOS)
OpenSSL/1.0.2k-fips PHP/7.4.9
mod_perl/2.0.11 Perl/v5.16.3
Last-Modified: Tue, 01 Mar 2016 18:57:50 GMT
ETag: "a5b-52d015789ee9e"
Accept-Ranges: bytes
Content-Length: 2651
Content-Type: text/html; charset=UTF-8
\r\n

HTTP response status codes

- status code appears in 1st line in server-to-client response message.
- some sample codes:

200 OK

- request succeeded, requested object later in this message

301 Moved Permanently

- requested object moved, new location specified later in this message (in Location: field)

400 Bad Request

- request msg not understood by server

404 Not Found

- requested document not found on this server

505 HTTP Version Not Supported

HTTP Recap

- ①. HTTP : client-server, request-response protocol
- ②. HTTP is stateless

Extending HTTP

- How to remember state?
- How to improve performance?
- How to scale?

HTTP is stateless but how does Instagram remember my login?

cookies are the mechanism
↓
Auxiliary to HTTP → still stateless]

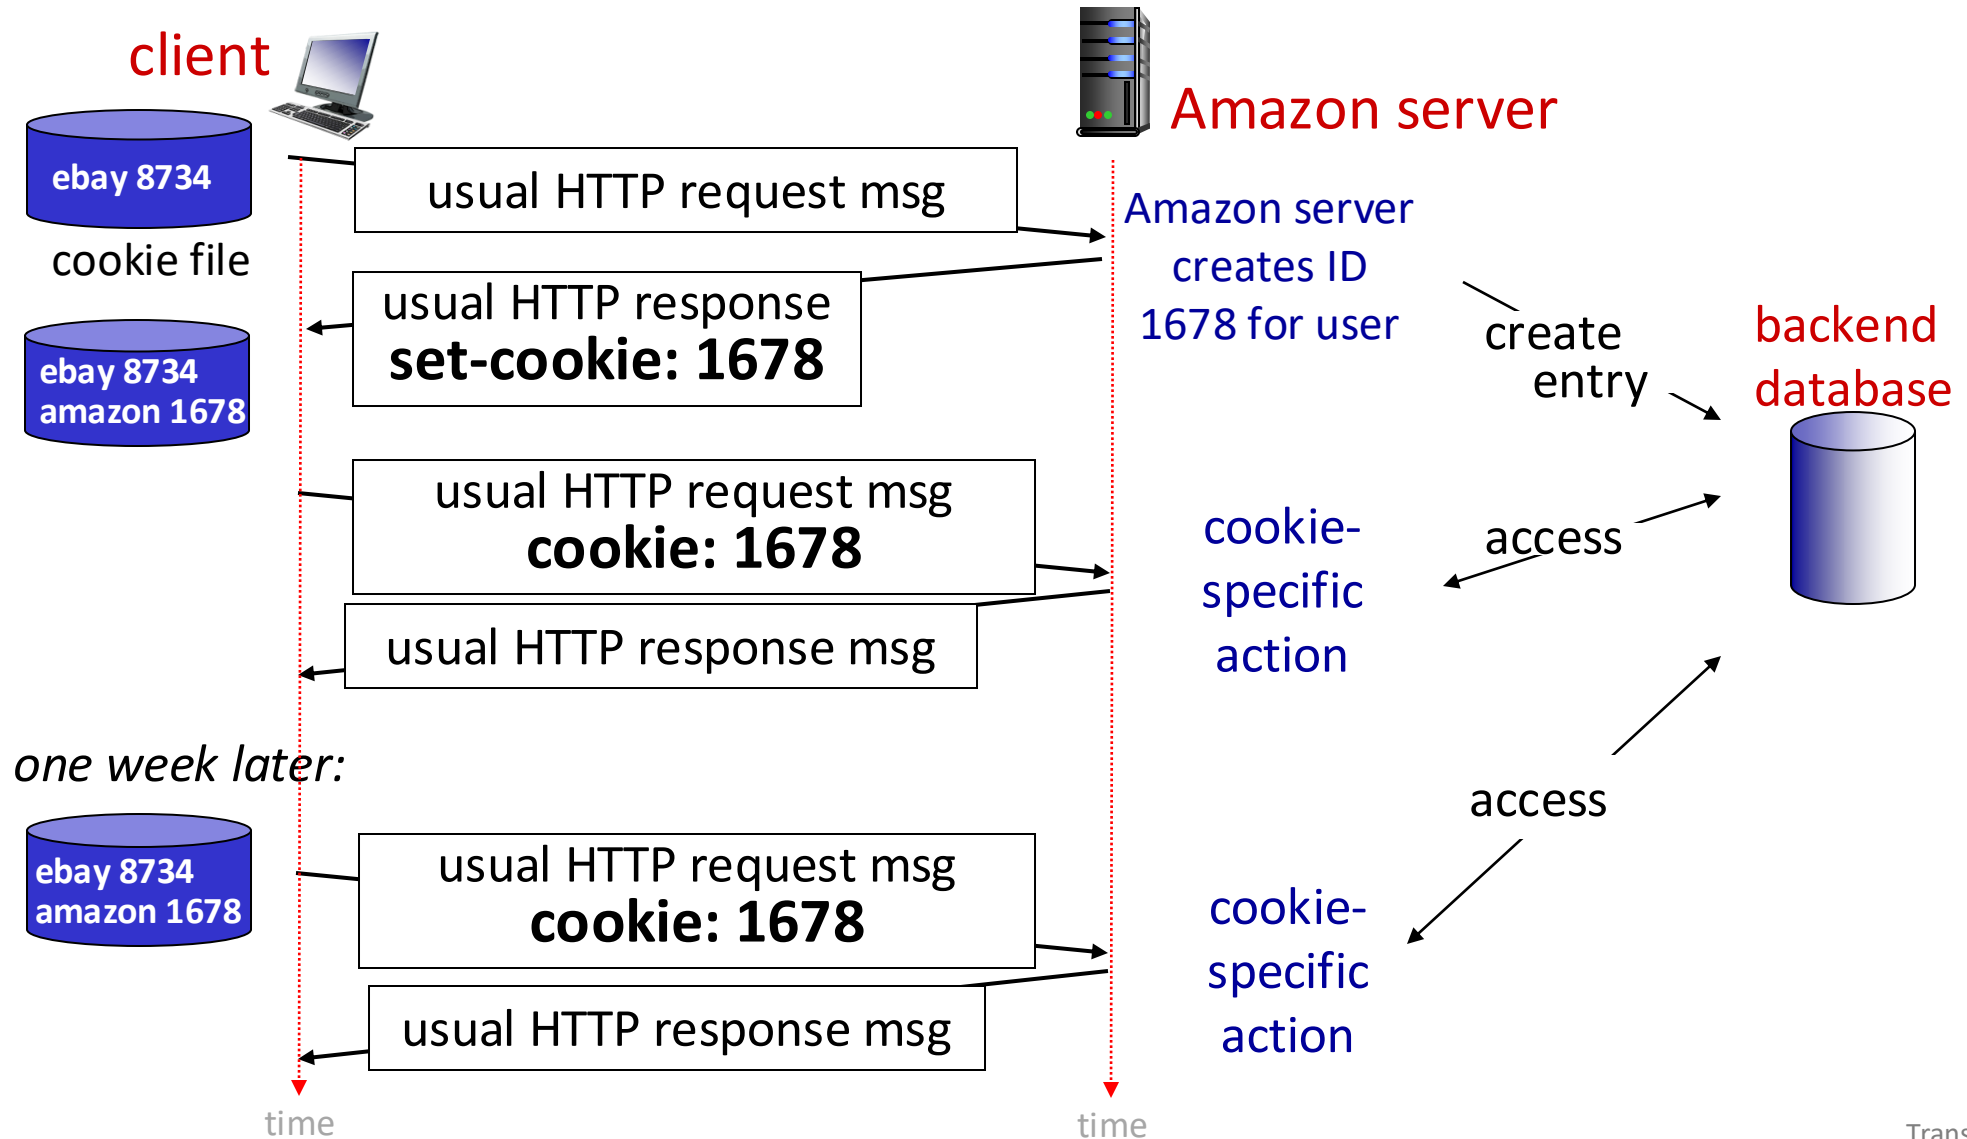
Maintaining user/server state: cookies

Web sites and client browser use *cookies* to maintain some state between transactions

four components:

- 1) cookie header line of HTTP *response* message
- 2) cookie header line in next HTTP *request* message
- 3) cookie file kept on user's host, managed by user's browser
- 4) back-end database at Web site

Maintaining user/server state: cookies




HTTP cookies: comments

What cookies can be used for:

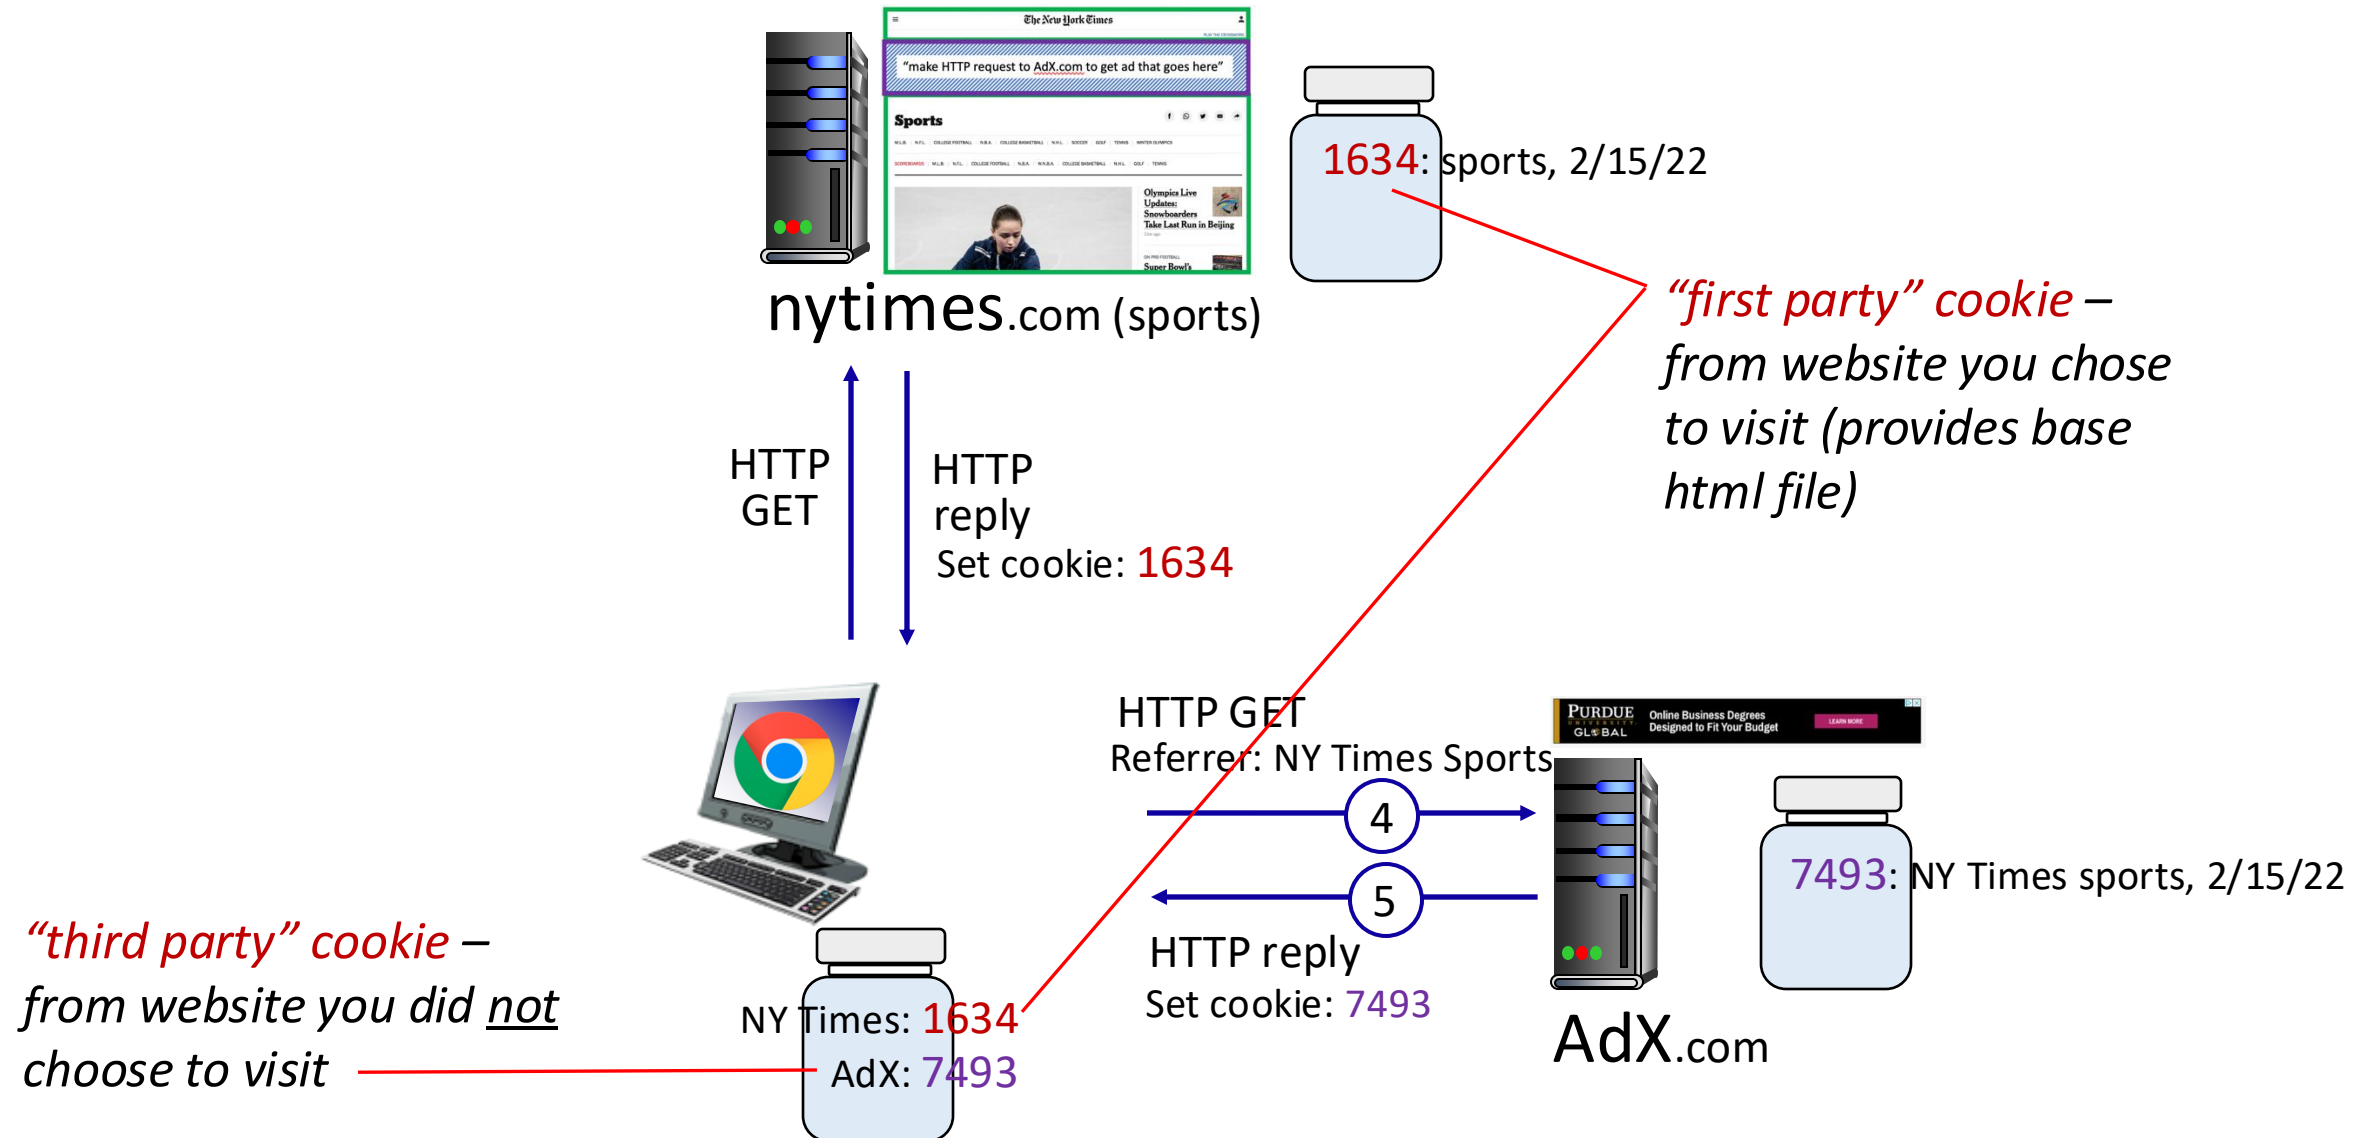
- authorization
- shopping carts
- recommendations (also ads)
- user session state (Web e-mail)

Example: displaying a NY Times web page

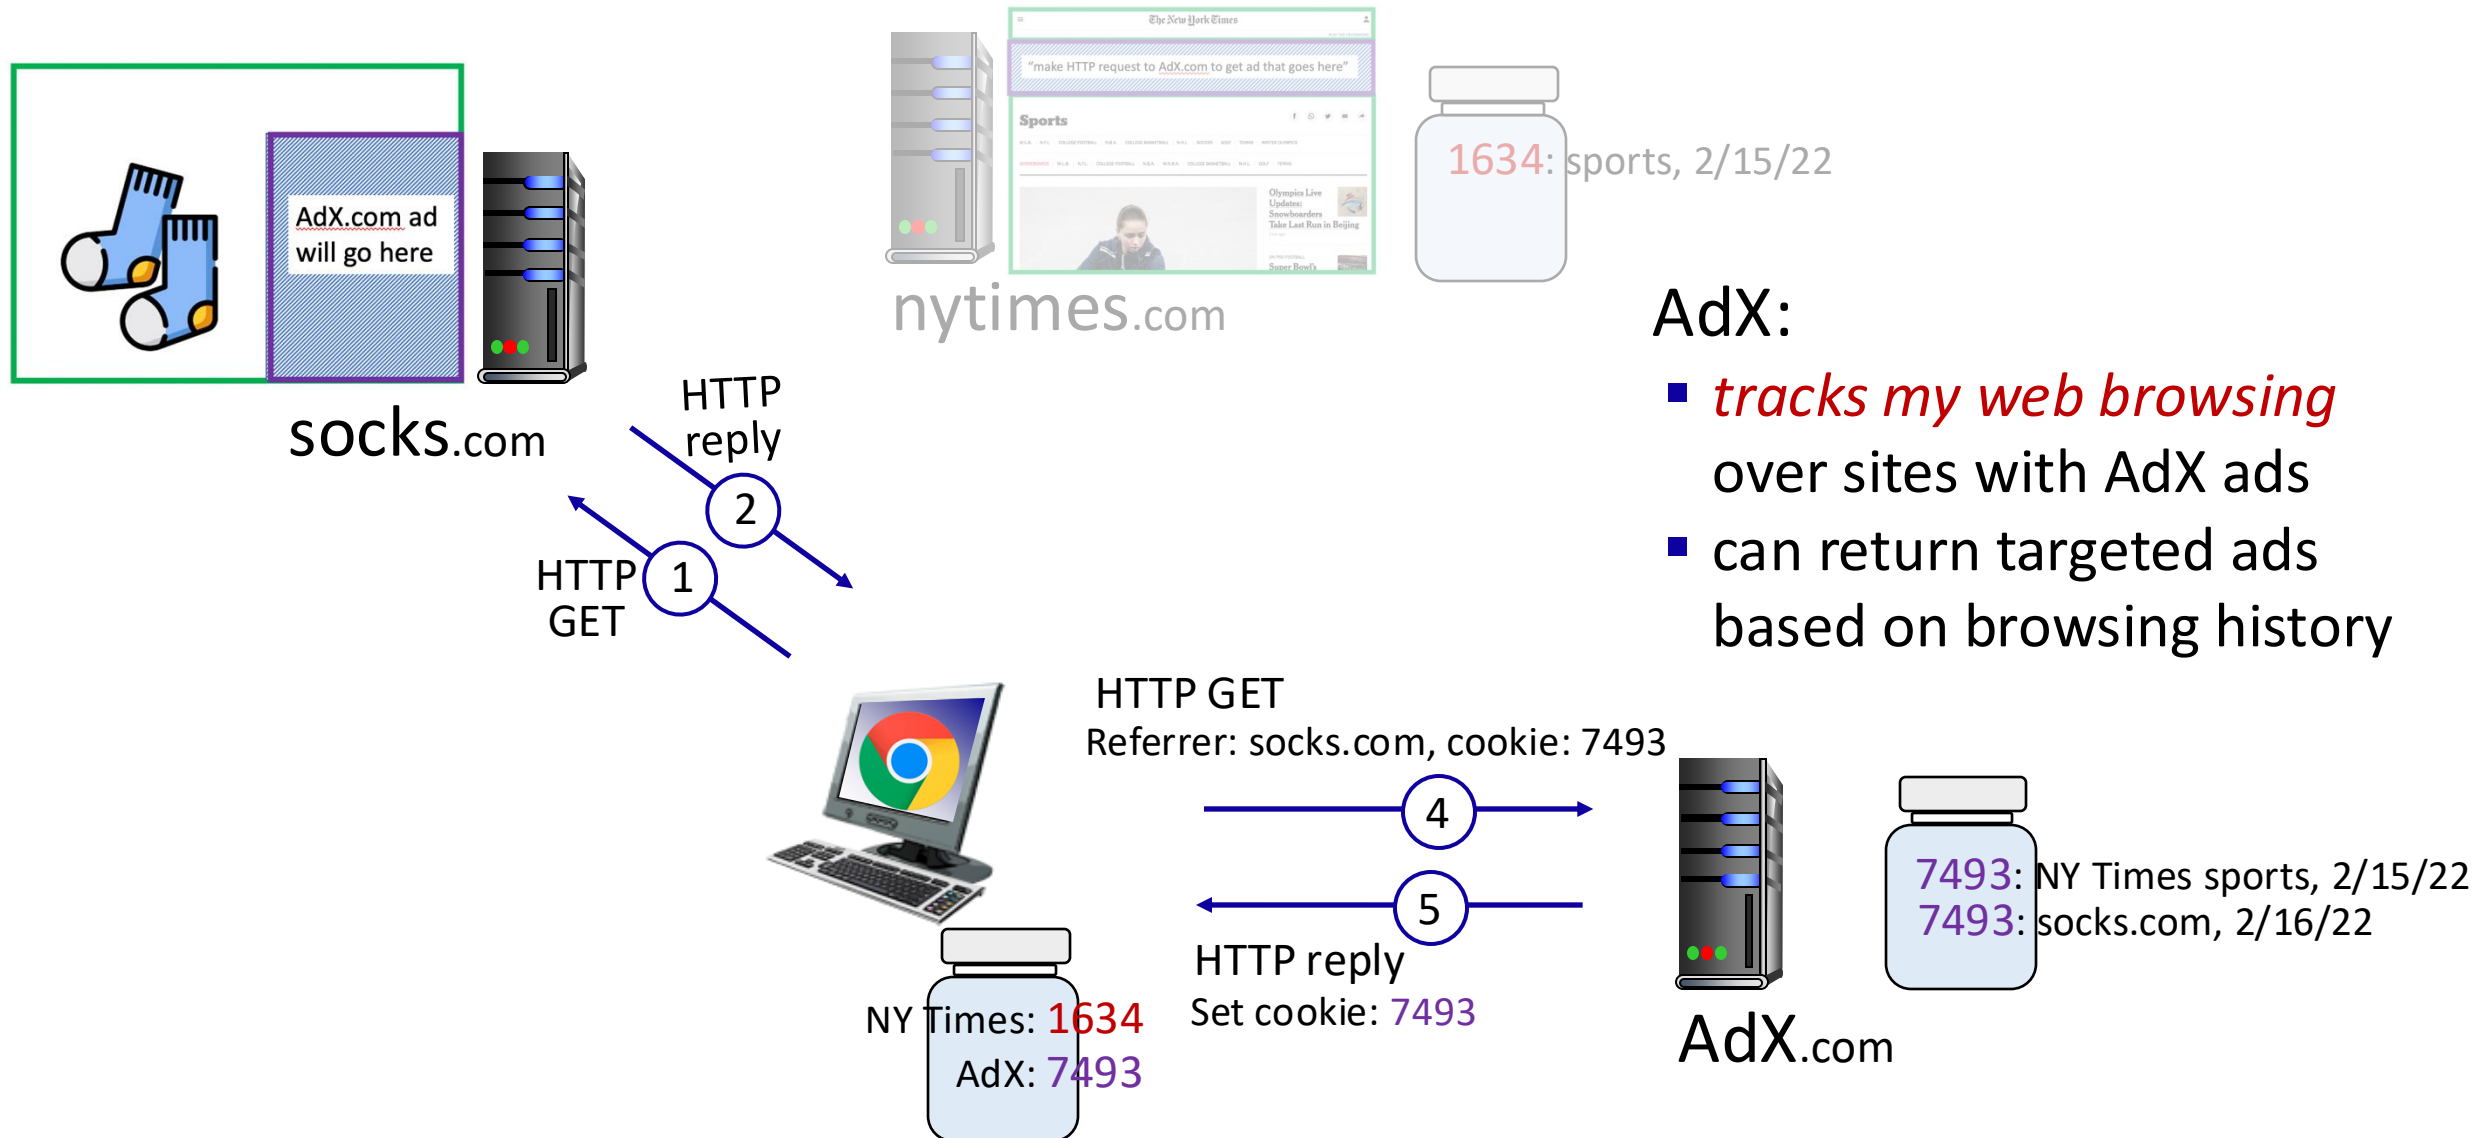
- ① GET base html file from nytimes.com
- ②
- ④ fetch ad from AdX.com
- ⑤
- ⑦ display composed page




Cookies: tracking a user's browsing behavior



Cookies: tracking a user's browsing behavior



Cookies: tracking a user's browsing behavior (one day later)



Cookies: tracking a user's browsing behavior

Cookies can be used to:

- track user behavior on a given website (**first party cookies**)
- track user behavior across multiple websites (**third party cookies**) without user ever choosing to visit tracker site (!)
- tracking may be *invisible* to user:
 - rather than displayed ad triggering HTTP GET to tracker, could be an invisible link

Data protection laws and cookies

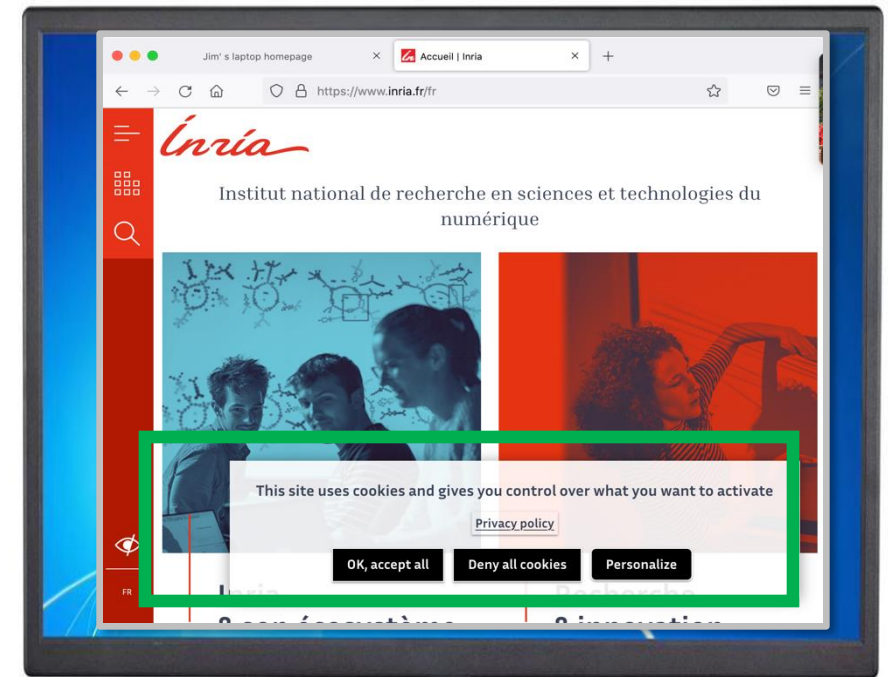
“Natural persons may be associated with online identifiers [...] such as internet protocol addresses, cookie identifiers or other identifiers [...].

This may leave traces which, in particular when combined with unique identifiers and other information received by the servers, may be used to create profiles of the natural persons and identify them.”

GDPR, recital 30 (May 2018)


India's data protection law: DPDPA

when cookies can identify an individual, cookies are considered personal data, subject to personal data regulations




User has explicit control over whether or not cookies are allowed

Improving HTTP



Serial download
(Inefficient)



HOL blocking
what if
obj 1 is
too big
↑
[Response sent] FCPS
request pipelining
(allowed in HTTP 1.1)

HTTP/2

Key goal: decreased delay in multi-object HTTP requests

HTTP1.1: introduced multiple, pipelined GETs over single TCP connection

- server responds *in-order* (FCFS: first-come-first-served scheduling) to GET requests
- with FCFS, small object may have to wait for transmission (**head-of-line (HOL) blocking**) behind large object(s)
- loss recovery (retransmitting lost TCP segments) stalls object transmission

HTTP/2

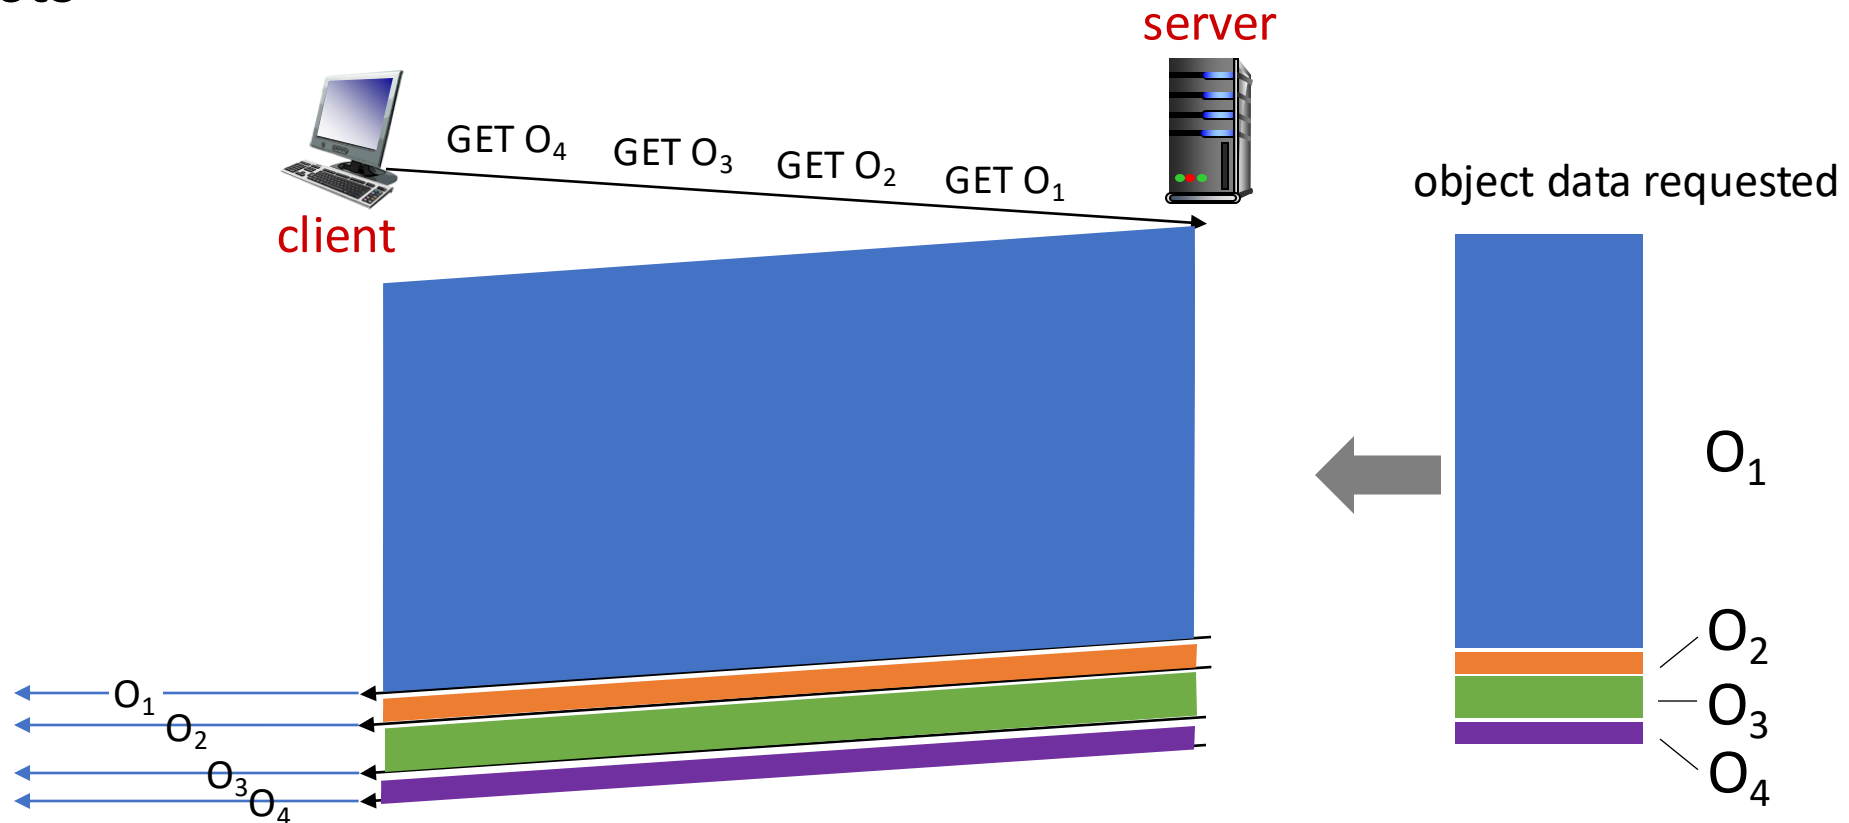
Key goal: decreased delay in multi-object HTTP requests

HTTP/2: [RFC 7540, 2015] increased flexibility at *server* in sending objects to client:

- methods, status codes, most header fields unchanged from HTTP 1.1
- transmission order of requested objects based on client-specified object priority (not necessarily FCFS)
- *push* unrequested objects to client
- divide objects into frames, schedule frames to mitigate HOL blocking

HTTP/2: mitigating HOL blocking

HTTP 1.1: client requests 1 large object (e.g., video file) and 3 smaller objects

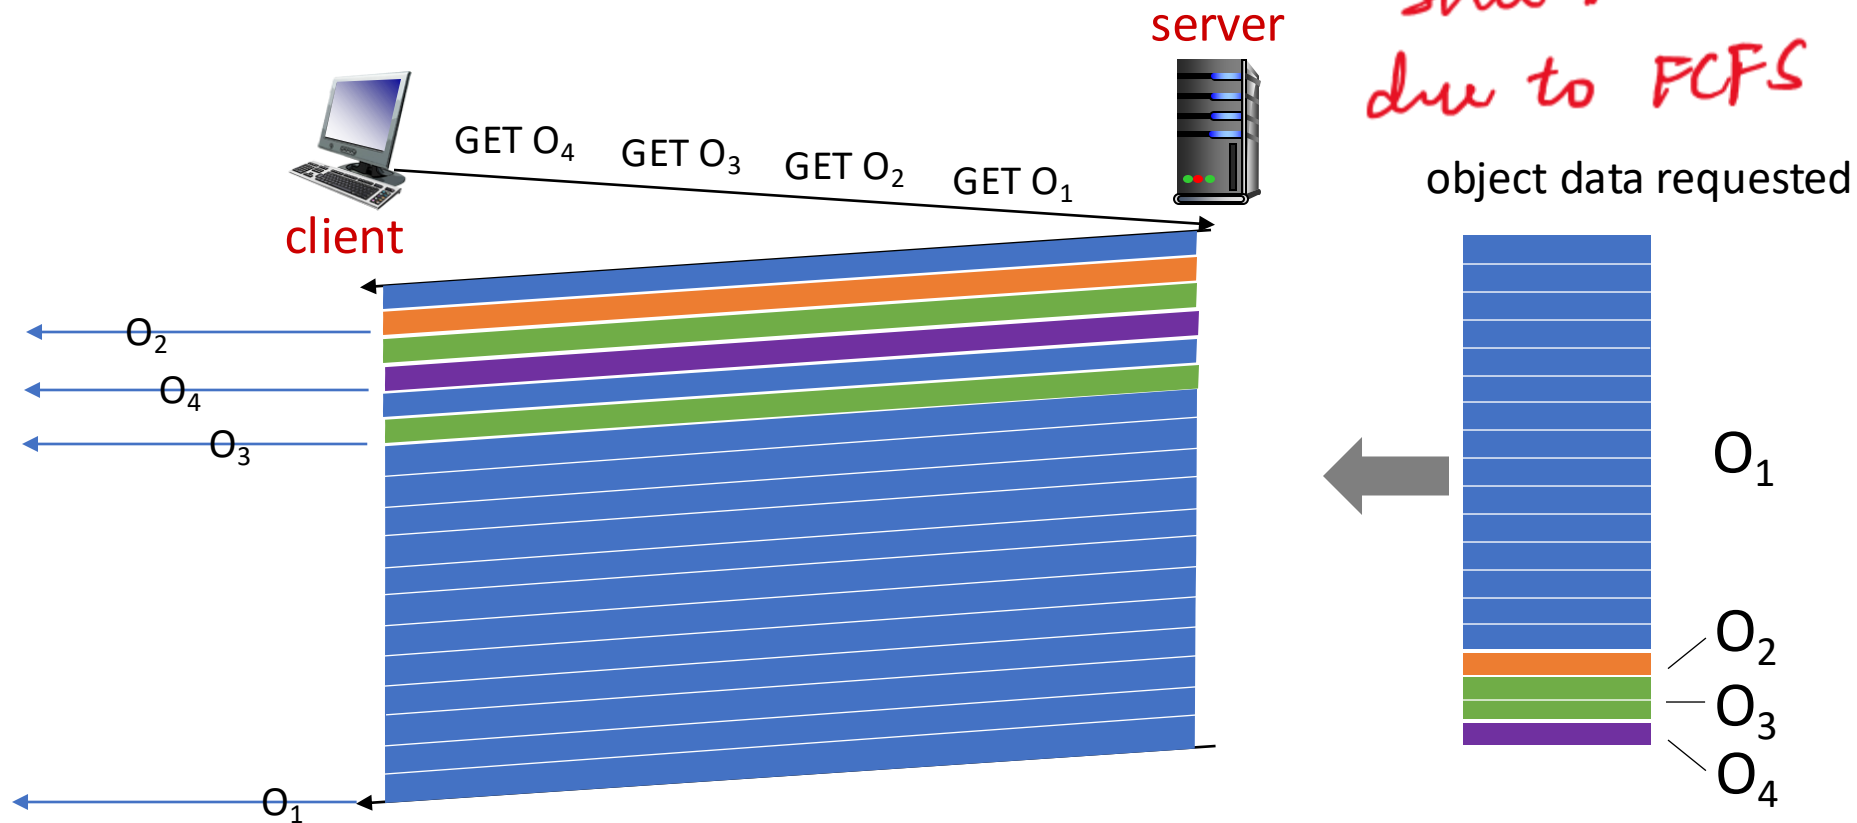


→ QUIC (or HTTP/3) over UDP [Mitigates TCP HOL blocking]

HTTP/2: mitigating HOL blocking

HTTP/2: objects divided into frames, frame transmission interleaved

still has HOL blocking due to FCFS in TCP



O₂, O₃, O₄ delivered quickly, O₁ slightly delayed

Improving HTTP


Goal: Scale content distribution

Caching: satisfy client requests without involving origin server



[A single server does
not scale]


Where in the network path can caching occur?



Browser caching: Conditional GET

Goal: don't send object if browser has up-to-date cached version

- no object transmission delay (or use of network resources)
- **client:** specify date of browser-cached copy in HTTP request
If-modified-since: <date>
- **server:** response contains no object if browser-cached copy is up-to-date:
HTTP/1.0 304 Not Modified




Content distribution networks (CDNs)

- CDN: geographically distribute collection of *server surrogates*
- Servers can be leased by many customers
- Popular CDNs: Limelight, Akamai, Level3
- Two kinds of server placement policies:
 - *enter deep*: push CDN servers deep into many access networks
 - close to users
 - Akamai: 240,000 servers deployed in > 120 countries (2015)
 - *bring home*: smaller number (10's) of larger clusters in POPs near access nets
 - used by Limelight



CDN architecture and cluster selection

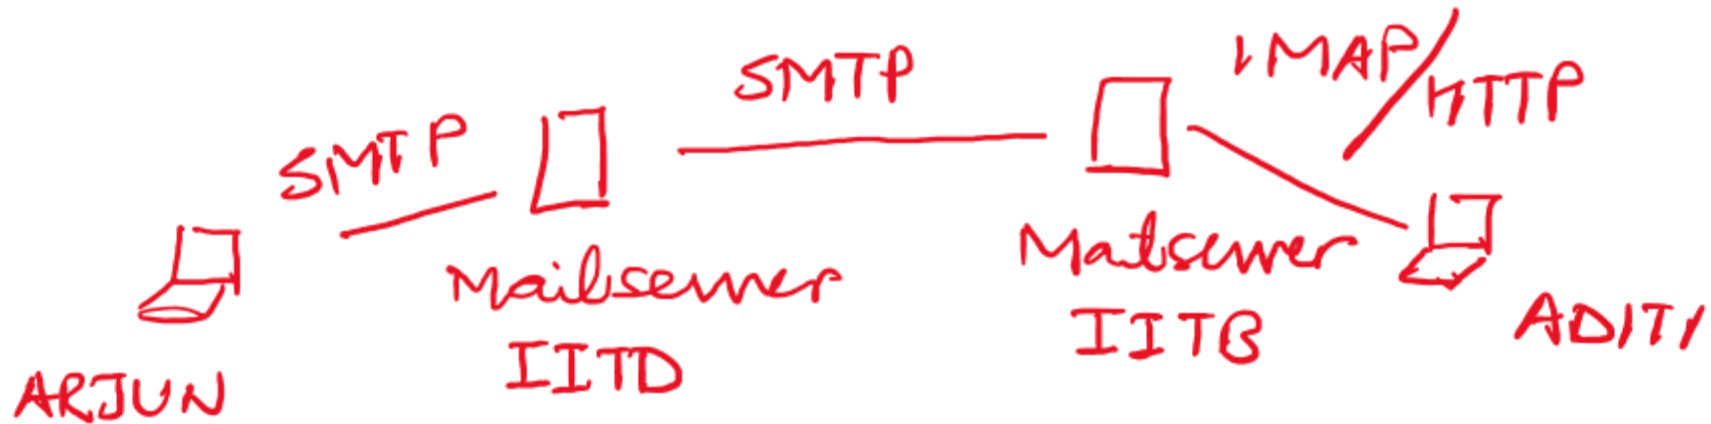


Policies and mechanisms for routing a client request to a cluster?

Journey of an Email

- Scenario: Arjun from IITD wants to send an email to his friend Aditi in IITB

What are the major components involved in email application?



Three major components:


- user agents
- mail servers
- simple mail transfer protocol: SMTP

Mail message format

SMTP: protocol for exchanging e-mail messages, defined in RFC 5321 (like RFC 7231 defines HTTP)

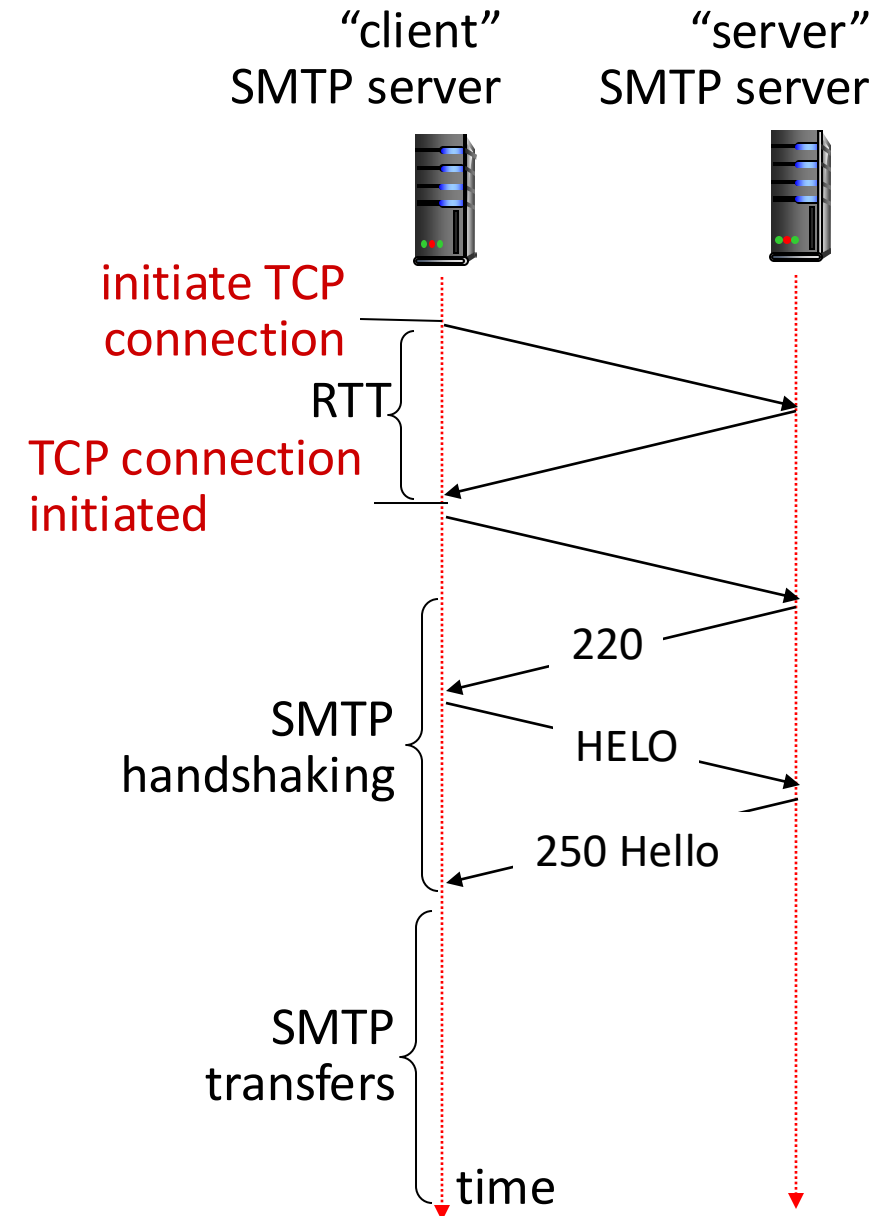
RFC 2822 defines *syntax* for e-mail message itself (like HTML defines syntax for web documents)

- header lines, e.g.,
 - To:
 - From:
 - Subject:these lines, within the body of the email message area different from SMTP MAIL FROM:, RCPT TO: commands!
- Body: the “message” , ASCII characters only



SMTP RFC (5321)

- uses TCP to reliably transfer email message from client (mail server initiating connection) to server, port 25
 - direct transfer: sending server (acting like client) to receiving server
- three phases of transfer
 - SMTP handshaking (greeting)
 - SMTP transfer of messages
 - SMTP closure
- command/response interaction (like HTTP)
 - **commands**: ASCII text
 - **response**: status code and phrase



Sample SMTP Interaction

S: 220 mail.iitb.ac.in

C: HELO mail.iitd.ac.in

S: 250 mail.iitb.ac.in Hello mail.iitd.ac.in, pleased to meet you

C: MAIL FROM:<user@iitd.ac.in>

S: 250 2.1.0 Sender OK

C: RCPT TO:<student@iitb.ac.in>

S: 250 2.1.5 Recipient OK

C: DATA

S: 354 Start mail input; end with <CRLF>.<CRLF>

C: Subject: Collaboration Request

C: From: user@iitd.ac.in

C: To: student@iitb.ac.in

C:

C: Hello, I would like to discuss a research collaboration.

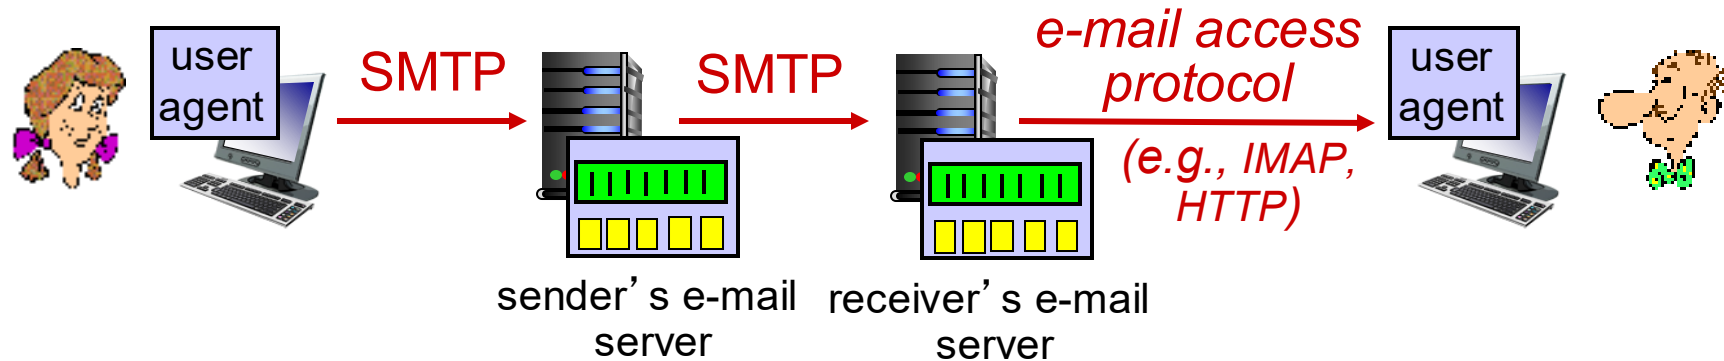
C: .

S: 250 2.0.0 Message accepted for delivery

C: QUIT

S: 221 2.0.0 mail.iitb.ac.in closing connection

Retrieving email: mail access protocols



- **SMTP:** delivery/storage of e-mail messages to receiver's server
- mail access protocol: retrieval from server
 - **IMAP:** Internet Mail Access Protocol [RFC 3501]: messages stored on server, IMAP provides retrieval, deletion, folders of stored messages on server
- **HTTP:** gmail, Hotmail, Yahoo!Mail, etc. provides web-based interface on top of SMTP (to send), IMAP (or POP) to retrieve e-mail messages

Configuring Mailbox

Outbox server (SMTP)

Username:
tmangla

Password:
●●●●●●●●●●

Server:
smtp.iitd.ac.in

Port:
587

Protection:
STARTTLS

Inbox server (IMAP)

Username:
tmangla

Password:
●●●●●●●●●●

Server:
mailstore.iitd.ac.in

Port:
993

Protection:
SSL

SMTP: observations

comparison with HTTP:

- HTTP: client pull
- SMTP: client push
- both have ASCII command/response interaction, status codes
- HTTP: each object encapsulated in its own response message
- SMTP: multiple objects sent in multipart message

SMTP/Email:

- Open standards
- Interoperability among email clients
- Exemplify the design spirit of Internet