

Computer Networks

COL 334/672

Application Layer: DNS

Sem 1, 2025-26

Recap: Application Layer

- HTTP
- Email
- **DNS** *or Domain Name System*
- P2P
- Video streaming

DNS: Domain Name System

DHCP : Application - layer protocol

- Humans understand names (google.com), *Name*
- Internet hosts, routers understand IP address (12.123.12.12) *Addressing*
- Q: how to map between IP address and name, and vice versa ?

Two questions:

- How to design the database?
- How to retrieve the IP for a given domain name?

Domain Name System (DNS):

- *phone book* of the Internet
- *application-layer protocol*: hosts, DNS servers communicate to *resolve* names (address/name translation)
 - *note*: core Internet function, **implemented as application-layer protocol**
 - complexity at network's "edge"
End-to-end design principle

How would you design the DNS database?

- **Primary goal:** a database of domain to IP mappings
- What are other design constraints/challenges?

{ a single machine
store all mappings }

↳ Scalable : # of queries to the table
& size of the table

↳ Fast (Performance)

↳ Secure

↳ Resilience

↳ Dynamic

And
outage

Approach 1: Centralized DNS

- single point of failure
- traffic volume
- maintenance
- ..

Replicate the database across
N servers

- ↳ ① scalability & Performance issues
② Expensive to maintain & update

Alternative

Split the database and store on multiple servers

Distributed and Hierarchical System

split the database?

■ Q: On what basis to decentralize?

Hierarchical domain name space

E.g. ~~cse~~.iitd.ac.in - IP mapping

① split is alphabetically

iitd.ac.in

cse.iitd.ac.in

→ could be
on the
same server


(Autonomy engrained into this)

Reverse DNS lookup

IP → hostname

② Split based on the

domain name structure




Decentralized and distributed system

(10.6.334.com, IP address)

- Q: On what basis to decentralize? Hierarchical domain name space
- Partition domain name hierarchy into zones managed by some authority → DNS registrars
 - E.g., ICANN is responsible for storing information about top-level domains
- Each zone corresponds to a name server

Name servers

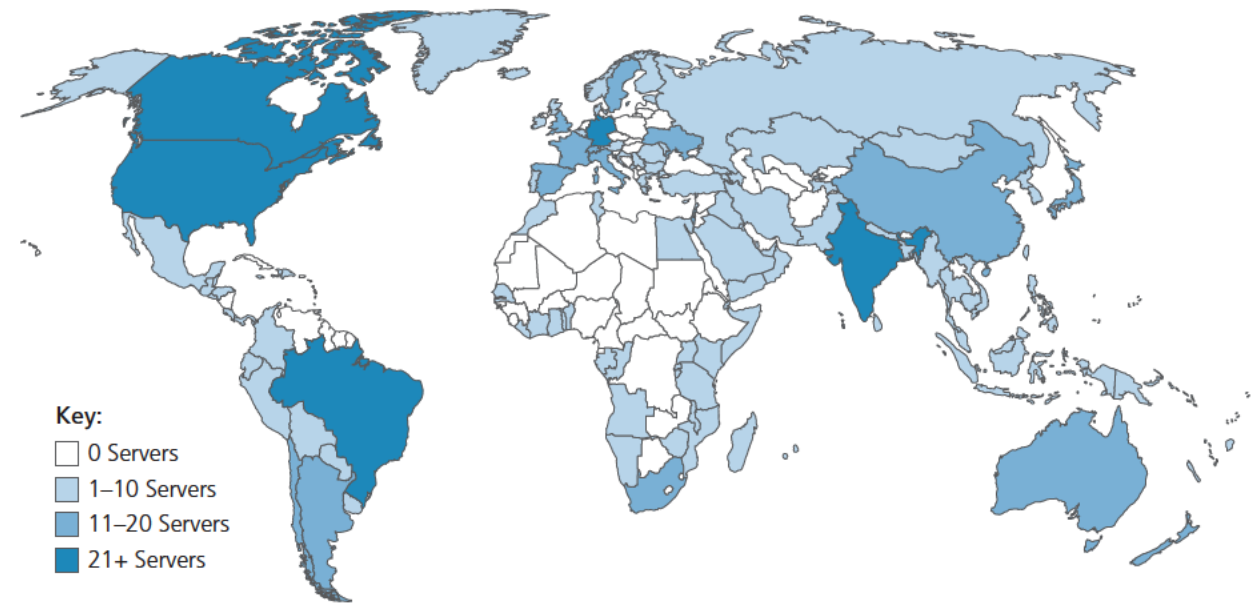


- Name servers are replicated and may be geographically distributed for reliability

DNS: root name servers

- official, contact-of-last-resort by name servers that can not resolve name
- *incredibly important* Internet function
 - Internet couldn't function without it!
 - DNSSEC – provides security (authentication, message integrity)
- ICANN (Internet Corporation for Assigned Names and Numbers) manages root DNS domain

13 logical root name “servers” worldwide each “server” replicated many times (~200 servers in US)



DNS records

DNS: distributed database storing resource records (RR)

RR format: (name, value, type, ttl) → *used for caching*

type=A

- name is hostname
- value is IP address

type=NS

- name is domain (e.g., foo.com)
- value is hostname of authoritative name server for this domain *ns.foo.com*

type=CNAME

- name is alias name for some “canonical” (the real) name
- www.ibm.com is really servereast.backup2.ibm.com
- value is canonical name

type=MX

- value is name of SMTP mail server associated with name

Example

AABA type record - IPv6

(Top-level domain)

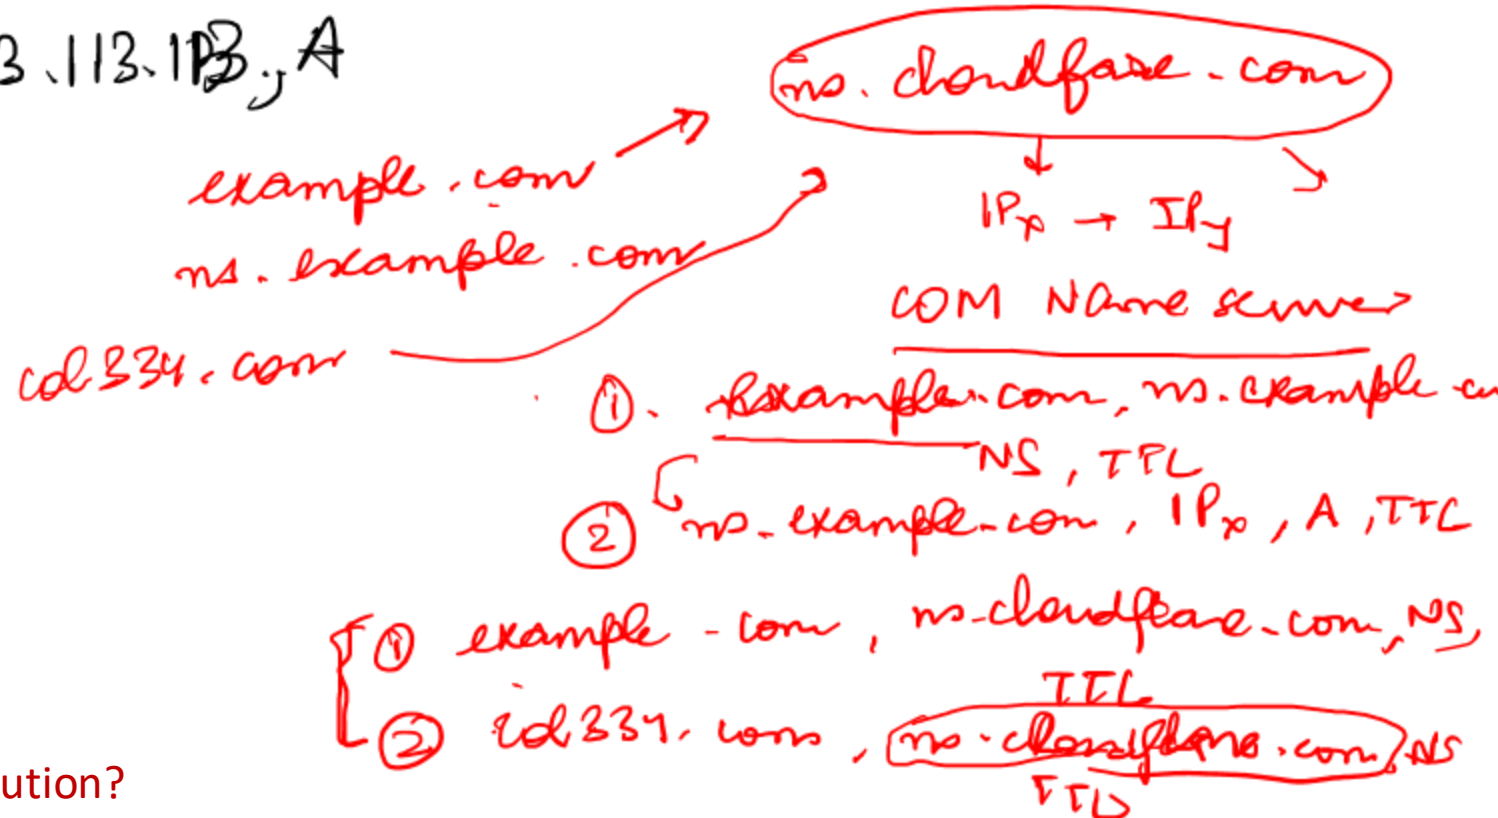
Root Name server

- ① .in, dns-xyz.in, NS, TTL
 - ② dns-xyz.in, 112.112.112.112, A,
 - ③ .com, dns-xyz.com, NS
 - ④ dns-xyz.com, 113.113.113.113, A
- edu
- net.

(who has authority)
where is it?

IN name server

- ① .ernet.in, dns-ernet.in, NS, TTL
- ② dns-ernet.in, x.y.z.w, A



How to do name resolution?

~~1800~~ 32 bit IP address

How To Do DNS Resolution?

UDP, TCP (TCP handshake)

- What is the transport protocol?

we use UDP (Reliability can be implemented in application layer)


Browser



- Does the browser directly query the root server?

↓
caching is very useful

Aggregation



DNS resolver

Local DNS name servers

→ DHCP → gateway IP
→ DNS resolver

Public DNS resolver: 8.8.8.8 → Google
1.1.1.1 → Cloudflare


- when host makes DNS query, it is sent to its *local* DNS server
 - Local DNS server returns reply, answering:
 - from its local cache of recent name-to-address translation pairs (possibly out of date!)
 - forwarding request into DNS hierarchy for resolution
 - each ISP has local DNS name server; to find yours:
 - MacOS: `% scutil --dns`
 - Windows: `>ipconfig /all`
- local DNS server doesn't strictly belong to hierarchy

DNS name resolution: iterated query

Example: host at iitd.ac.in wants IP address for google.com

Iterated query:

- contacted server replies with name of server to contact
- “I don’t know this name, but ask this server”



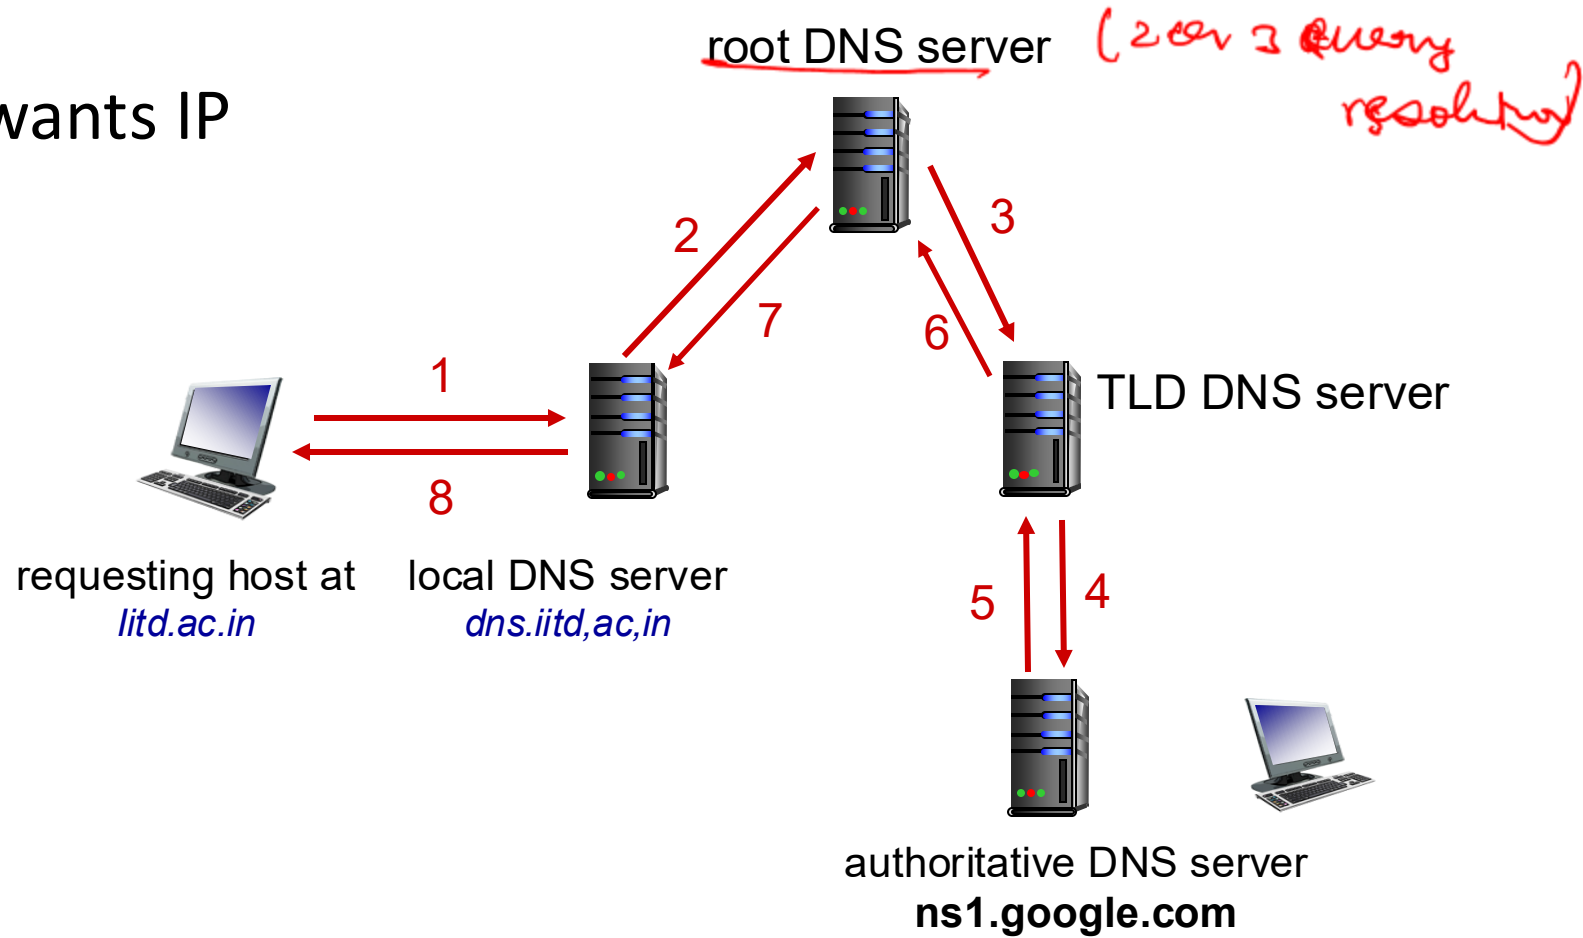
iitd-ac-in

DNS name resolution: recursive query

Example: host at iitd.ac.in wants IP address for google.com

Recursive query:

- puts burden of name resolution on contacted name server
- heavy load at upper levels of hierarchy?



DNS protocol messages

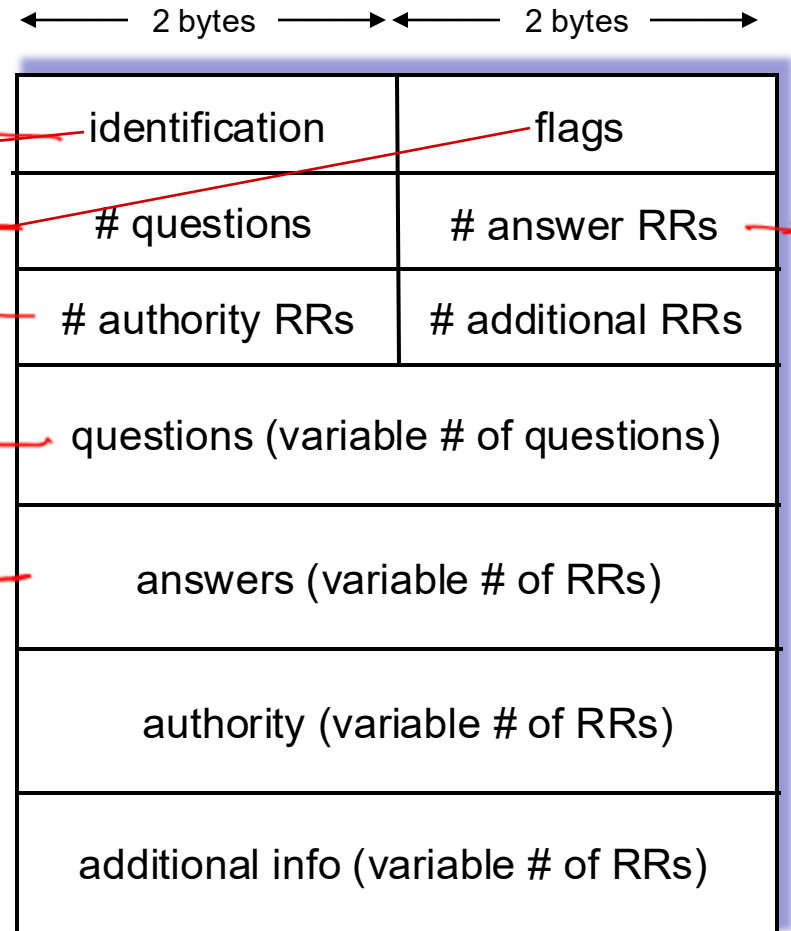
DNS *query* and *reply* messages, both have same *format*:

message header:

- **identification**: 16 bit # for query, reply to query uses same #

it's ■ **flags**:

- query or reply
- recursion desired
- recursion available
- reply is authoritative

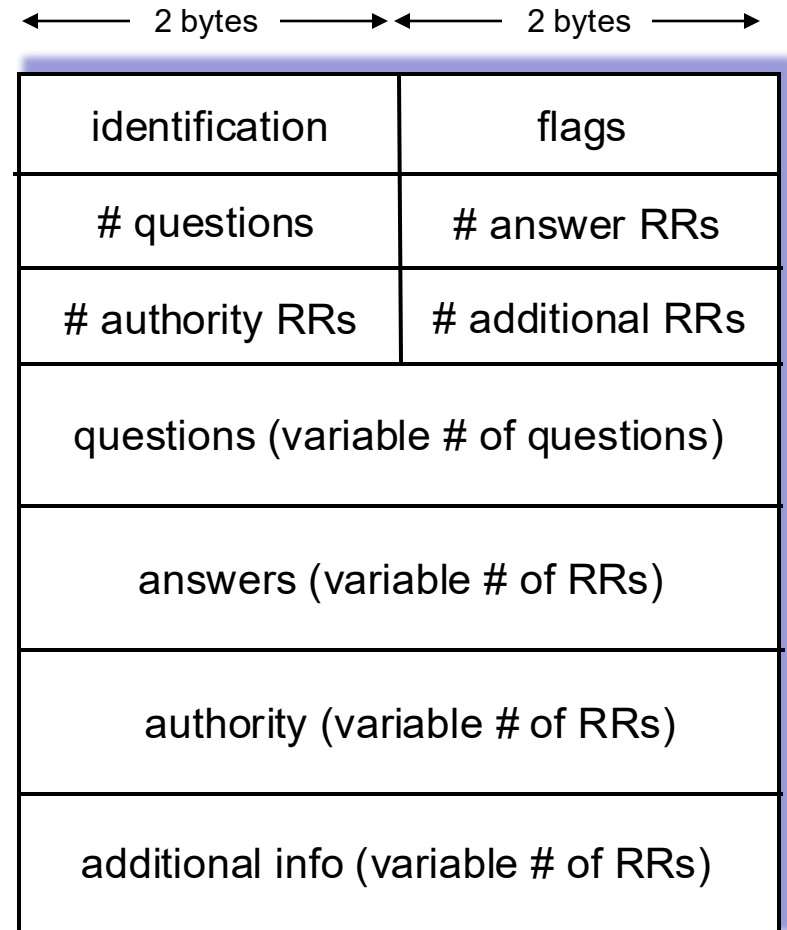


DNS protocol messages

DNS *query* and *reply* messages, both have same *format*:

dig google.com

```
✓ Queries
  ✓ google.com: type A, class IN
    Name: google.com
    [Name Length: 10]
    [Label Count: 2]
    Type: A (Host Address) (1)
    Class: IN (0x0001)
✓ Answers
  ✓ google.com: type A, class IN, addr 142.250.194.142
    Name: google.com
    Type: A (Host Address) (1)
    Class: IN (0x0001)
    Time to live: 227 (3 minutes, 47 seconds)
    Data length: 4
    Address: 142.250.194.142
```



Caching DNS Information

- once (any) name server learns mapping, it *cached* mapping, and *immediately* returns a cached mapping in response to a query
 - caching improves response time
 - cache entries timeout (disappear) after some time (TTL)
 - TLD servers typically cached in local name servers
- cached entries may be *out-of-date*
 - if named host changes IP address, may not be known Internet-wide until all TTLs expire!
 - *best-effort name-to-address translation!*

Getting your info into the DNS

example: new startup “Network Utopia”

- register name networkutopia.com at DNS registrar (e.g., Network Solutions)
 - provide names, IP addresses of authoritative name server (primary and secondary)
 - registrar inserts NS, A RRs into .com TLD server:
 - [(networkutopia.com, dns1.networkutopia.com, NS)
 - (dns1.networkutopia.com, 212.212.212.1, A)
- create authoritative server locally with IP address 212.212.212.1
 - type A record for www.networkutopia.com
 - type MX record for networkutopia.com

DNS observations

xyz. github.is Nameserver likely
↳ Managed by Github

DDoS attacks

- bombard root servers with traffic
 - not successful to date
 - traffic filtering
 - local DNS servers cache IPs of TLD servers, allowing root server bypass

Spoofing attacks

- intercept DNS queries, returning bogus replies
 - DNS cache poisoning
 - RFC 4033: DNSSEC authentication services

Centralization of DNS

- Name servers hosted by third-party (e.g., cloudflare, amazon)
 - Why?
 - Single point of failure?