Computer Networks
COL 334/672

Application Layer: P2P

Slides adapted from KR

Sem 1, 2025-26

DNS and Load Balancing

IP Geolocation of the resolver

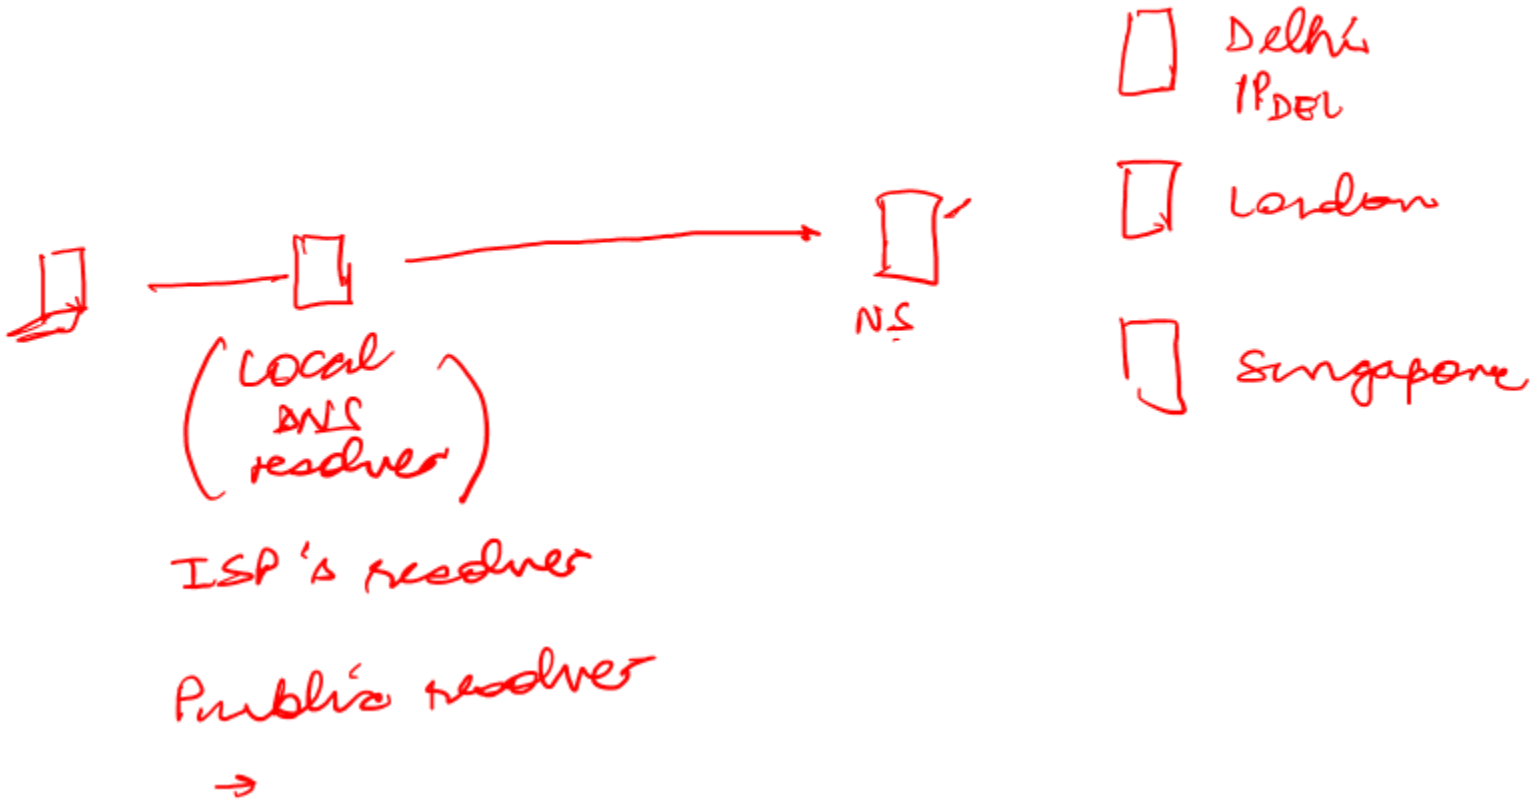

DNS, Content Regulation, Censorship

use VPN or public DNS resolver

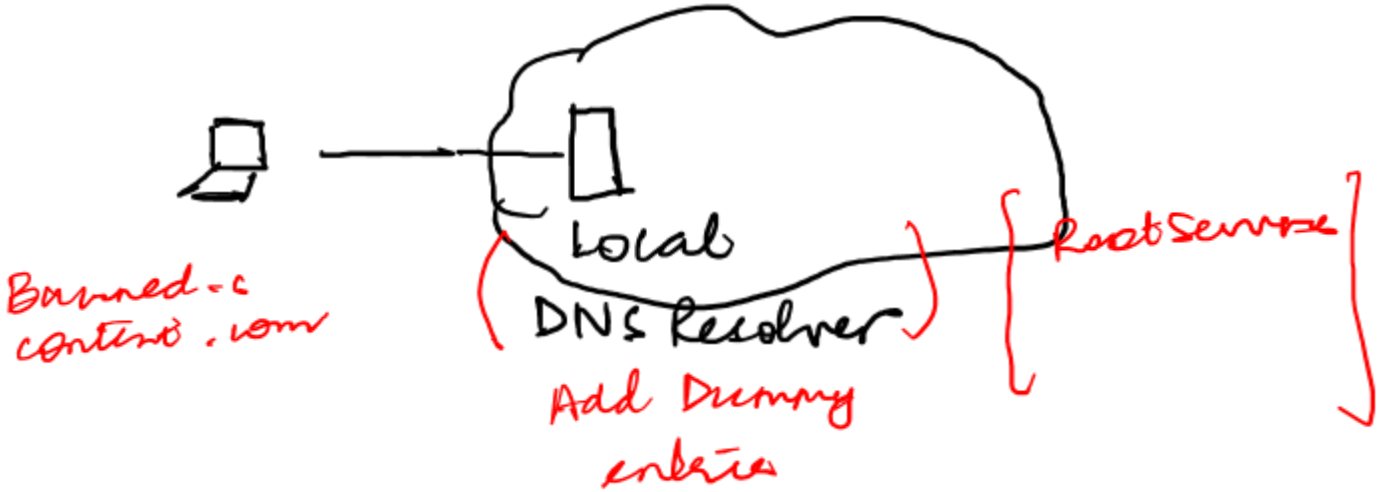

Google DNS resolver

Recap: Application Layer

- HTTP
- Email
- DNS
- **P2P**
- Video streaming

What is Peer to Peer (P2P) Communication?

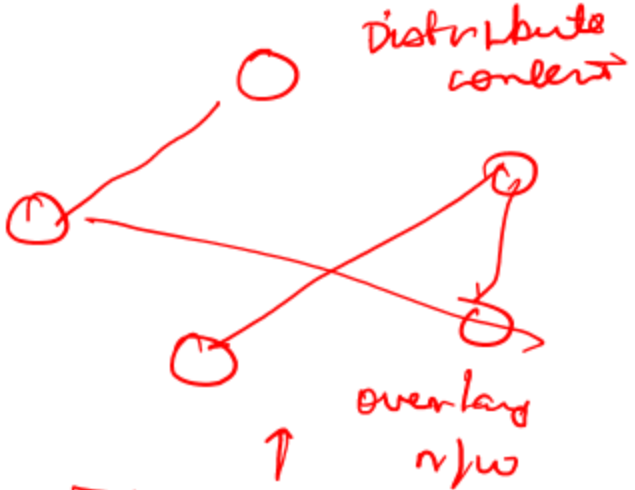

{ Whatsapp: Voice / Video Telephony (

Which applications?

Torrent : Content Distribution

Distribute content using this P2P n/w

Why P2P for content distribution?

→ PROS: self scalable

CONS: security issue, Peers can come & go

- Scales better than client-server architecture **Why?**

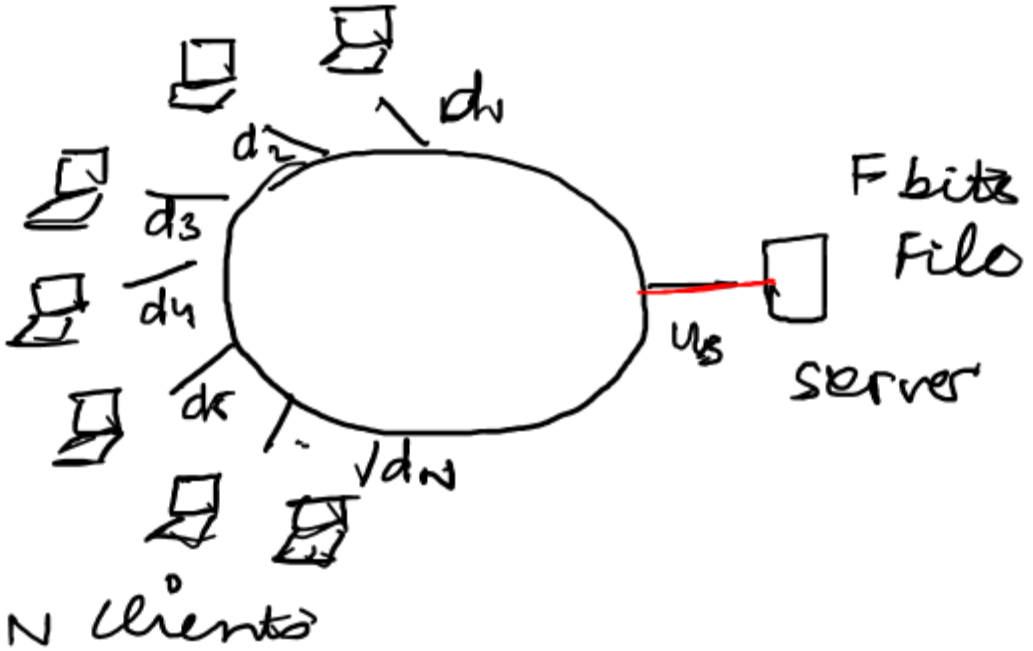

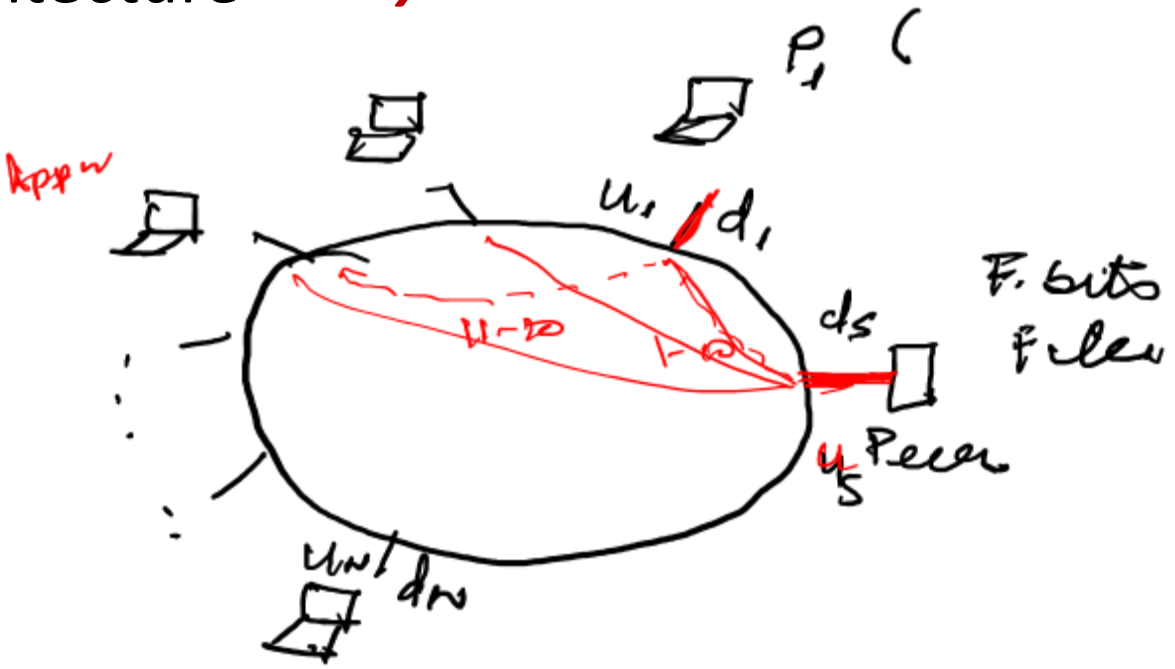

$$\max\left\{\frac{NF}{u_s}, \frac{F}{d_{min}}\right\}$$

as N scales, this does not scale

[unless you have more servers]

$$\max\left\{\frac{F}{u_s}, \frac{F}{d_{min}}, \frac{NF}{\sum u_i + u_s}\right\}$$

Particularly popular in the early 2000s

Two Interesting Questions for Content Distribution

- How to find content?
- How to download content?

How to Find Content?

- **Scenario:**
  - Network of peers who may each have a different set of content
  - Each peer is connected directly to some other peers, not necessarily all (why?)
- **Problem Statement**: How does a peer find who has file F_i?

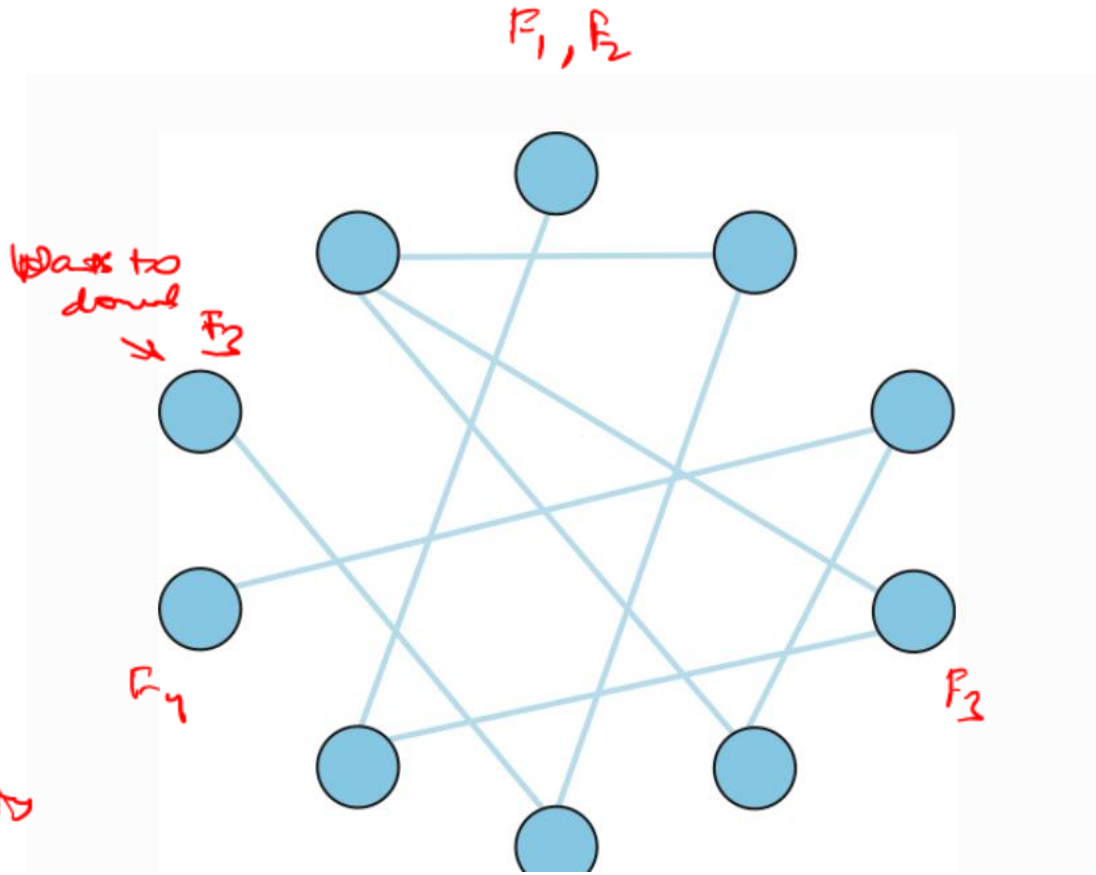

Broadcast → TTL
Seq #s

unnecessary overhead → latency

Centralized index table

→ Scalability
→ SPF
→ accountability

Finding a file in a P2P network

- Intuition: Some indexing is useful for a faster lookup.
- Challenge: But can't have a centralized hash table
- Solution: Use a distributed hash table

Distribute the hash table

(DHT)

(PASTRY) → DHT

Idea: H(filename) H(Name a IP)

- Map the objects and the nodes to a common virtual space
- Store the object information in a node that is closest to it in the virtual space

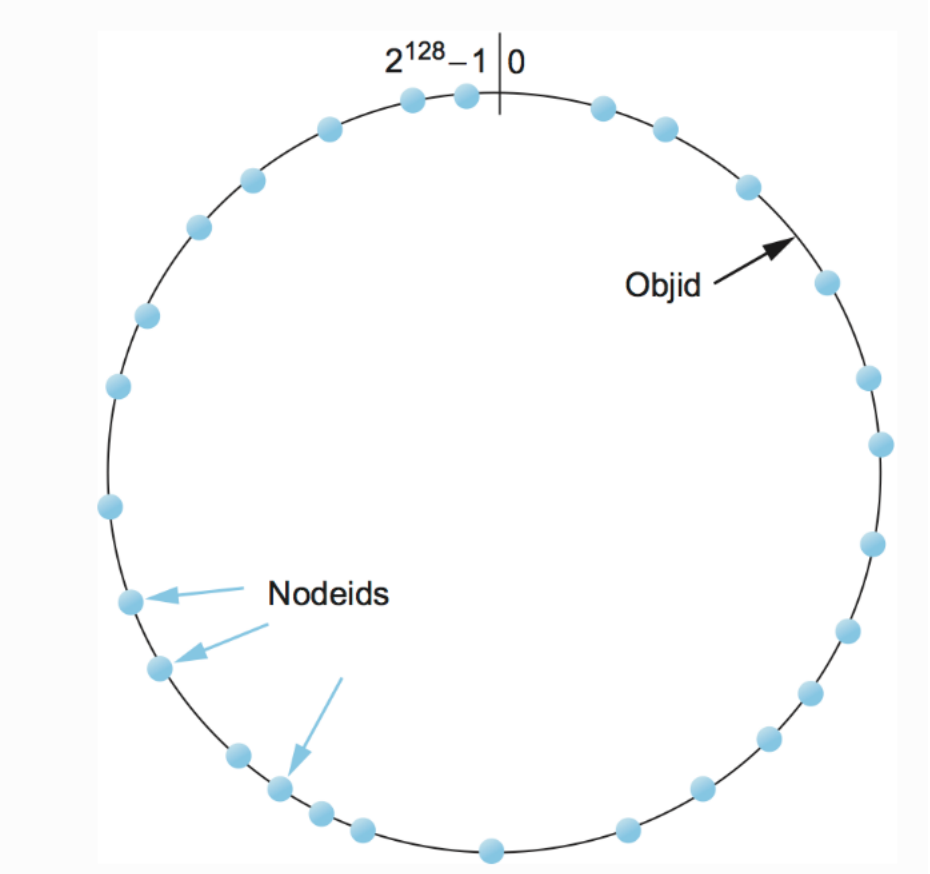

PASTRY Example

{0, ... 7}

Node : 2, 4, 6

Files : F_1 F_2 F_3
2, 3, 7

If node with same values as $H(F_i)$ present

Store the $\langle k, v \rangle$ in that node.

Else

store in successor $(H(F_i))$

Distibuted Hash Table

PASTRY: Searching Closest Node

Idea: *To search for a file f, route query messages closer to H(f) in the virtual space until you find the node containing information about f*

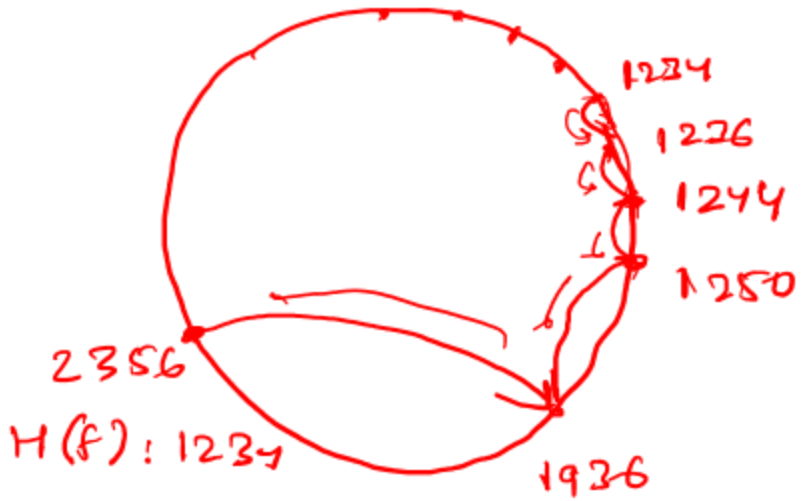

Challenge: *How do we ensure that we can always go to a closer node?*

Put some structure on the network

PASTRY: Searching Closest Node

CHORD : Puts more structure

Randomized algorithm

Idea: *To search for a file f, route query messages closer to H(f) in the virtual space until you find the node containing information about f*

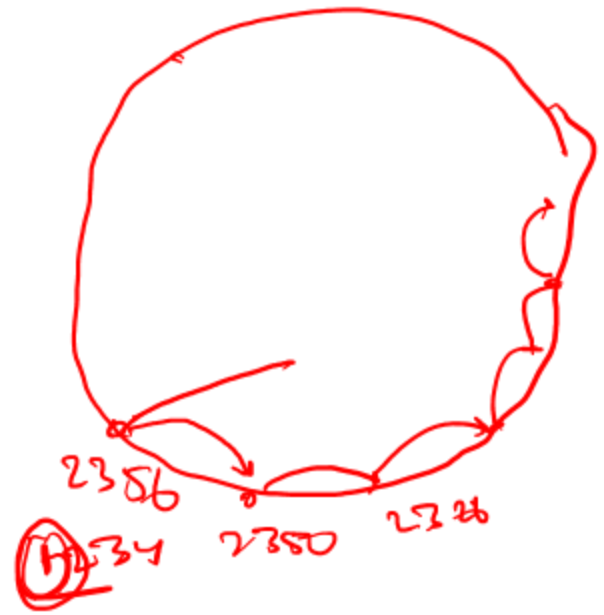

Challenge: *How do we ensure that we can always go to a closer node?*

Reach the neighbor in log(N)

Solution: *Each node should store L nodes (L/2 successors, L/2 predecessors) and log(N) nodes distributed randomly in the virtual space*

Physically closer

just linear search is too slow

let's make bigger jumps

Distributed Hash Table

- You should think about the following:
  - How the neighbors are maintained in the first place
  - What happens when a neighbor disconnects?
  - How is a new neighbor added?

- Various optimizations exist for DHTs

- Used in other domains such as distributed file system, web caching etc.

Next question: How to download content?